

# Homework 5

**name:** Vasudha Rohatgi

## INSTRUCTIONS:

- Submit the solutions to all questions as a single PDF file on Gradescope.
- Please show all your work to receive full credit.
- All answers must be typeset preferably using LaTeX (the template is provided, so please use it). Start each exercise on a new page and write your answers inside the `exercise` environments provided.
- For all coding questions, or questions for which you rely on any computations please include the code and the outputs that reproduce your analyses in your PDF (you can use Ctrl+P on your Colab notebook, save as PDF and merge with your submission document). Don't forget to run all cells before submitting.  
If a question only involves coding, write "Code provided below." in the solutions box and insert the PDF pages with the relevant code on the next page.
- When submitting on Gradescope, please indicate which pages of your solution are associated with each problem.
- You are free to discuss homework solutions with other students, but write-ups must be done independently. Be sure to give credit to any students you collaborate with.

Failure to comply with these instructions will result in a deduction of points.

## Optional Questions

1. Do you have any confusion or questions about the previous lectures?
2. Any suggestions or thoughts about the course?

## Suggestions

## Exercise 1

**Traffic Study on Residential Streets [30 points].** A survey was done of bicycle and other vehicular traffic in the neighborhood of the campus of the University of California, Berkeley, in the spring of 1993. Sixty city blocks were selected at random; each block was observed for one hour, and the numbers of bicycles and other vehicles traveling along that block were recorded. The sampling was stratified into six types of city blocks: busy, fairly busy, and residential streets, with and without bike routes, with ten blocks measured in each stratum. The following table displays the number of bicycles and other vehicles recorded in the study.

Type of street	Bike route?	Counts of bicycles/other vehicles
Residential	yes	16/58, 9/90, 10/48, 13/57, 19/103, 20/57, 18/86, 17/112, 35/273, 55/64
Residential	no	12/113, 1/18, 2/14, 4/44, 9/208, 7/67, 9/29, 8/154
Fairly busy	yes	8/29, 35/415, 31/425, 19/42, 38/180, 47/675, 44/620, 44/437, 29/47, 18/462
Fairly busy	no	10/557, 43/1258, 5/499, 14/601, 58/1163, 15/700, 0/90, 47/1093, 51/1459, 32/1086
Busy	yes	60/1545, 51/1499, 58/1598, 59/503, 53/407, 68/1494, 68/1558, 60/1706, 71/476, 63/752
Busy	no	8/1248, 9/1246, 6/1596, 9/1765, 19/1290, 61/2498, 31/2346, 75/3101, 14/1918, 25/2318

Now restrict your attention to the first four rows of the table: the data on residential streets.

- [8 points] Let  $y_1, \dots, y_{10}$  and  $z_1, \dots, z_8$  be the observed proportion of traffic that was on bicycles in the residential streets with bike lanes and with no bike lanes, respectively (so  $y_1 = 16/(16 + 58)$  and  $z_1 = 12/(12 + 13)$  for example). Set up a model so that the  $y_i$ 's are independent and identically distributed given parameters  $\theta_y$  and the  $z_i$ 's are independent and identically distributed given parameters  $\theta_z$ .
- [4 points] Set up a prior distribution that is independent in  $\theta_y$  and  $\theta_z$ .
- [10 points] Draw 1000 simulations from the posterior distribution. (Hint:  $\theta_y$  and  $\theta_z$  are independent in the posterior distribution, so they can be simulated independently.)
- [8 points] Let  $\mu_y = E(y_i|\theta_y)$  be the mean of the distribution of the  $y_i$ 's;  $\mu_y$  will be a function of  $\theta_y$ . Similarly, define  $\mu_z$ . Using your posterior simulations from 3, plot a histogram of the posterior simulations of  $\mu_y - \mu_z$ , the expected difference in proportions in bicycle traffic on residential streets with and without bike lanes.

### Solution

Part 1: Model Setup Let  $y_1, \dots, y_{10}$  and  $z_1, \dots, z_8$  be the observed proportions of traffic that were bicycles on residential streets with and without bike lanes, respectively. The proportions are calculated as:

$$y_1 = \frac{16}{16 + 58}, \quad z_1 = \frac{12}{12 + 13}, \quad \text{and so on.}$$

Assume the observed proportions,  $y_i$  and  $z_i$ , are independent and identically distributed given parameters  $\theta_y$  and  $\theta_z$ . Using a Beta distribution as a prior for the proportions, the likelihood is:

$$y_i | \theta_y \sim \text{Beta}(\alpha_y, \beta_y), \quad z_i | \theta_z \sim \text{Beta}(\alpha_z, \beta_z).$$

## Part 2: Prior Distribution

For simplicity, I chose independent uniform priors for  $\theta_y$  and  $\theta_z$ :

$$\theta_y \sim \text{Beta}(1, 1), \quad \theta_z \sim \text{Beta}(1, 1).$$

## Part 3: Posterior Simulation

Using the Beta-binomial conjugacy property, the posterior distributions are:

$$\theta_y | y \sim \text{Beta}(\alpha_y + \sum y_i n_i, \beta_y + \sum (1 - y_i) n_i),$$

$$\theta_z | z \sim \text{Beta}(\alpha_z + \sum z_i n_i, \beta_z + \sum (1 - z_i) n_i),$$

where  $n_i$  is the total number of vehicles observed for each block.

*The R code for posterior simulation is attached*

## Part 4: Expected Difference in Proportions

Define  $\mu_y = \mathbb{E}[y_i | \theta_y]$  and  $\mu_z = \mathbb{E}[z_i | \theta_z]$ . From the posterior distributions,  $\mu_y$  and  $\mu_z$  can be approximated as the means of the sampled posterior distributions for  $\theta_y$  and  $\theta_z$ , respectively.

The expected difference is computed as:

$$\mu_y - \mu_z = \frac{\alpha_y + \sum y_i n_i}{\alpha_y + \beta_y + \sum n_i} - \frac{\alpha_z + \sum z_i n_i}{\alpha_z + \beta_z + \sum n_i}.$$

## Results

Using the R code provided, I obtained the following posterior estimates:

- Posterior mean of  $\theta_y$ : `mean(theta_y_samples)`
- Posterior mean of  $\theta_z$ : `mean(theta_z_samples)`
- Posterior mean of  $\theta_y - \theta_z$ : `mean(diff_samples)`

## Conclusion

This Bayesian analysis provides insights into the expected difference in proportions of bicycle traffic on residential streets with and without bike lanes. The histogram of posterior differences is included in the output.

# STATS 551 - HW 5 - Q1

```

# Data
y <- c(16/74, 9/99, 10/58, 13/70, 19/121, 20/77, 18/104, 17/129, 35/308, 55/119)
z <- c(12/25, 1/19, 2/16, 4/48, 9/217, 7/74, 9/38, 8/162)

# Total counts
n_y <- c(74, 99, 58, 70, 121, 77, 104, 129, 308, 119)
n_z <- c(25, 19, 16, 48, 217, 74, 38, 162)

# Prior parameters
alpha_y <- 1; beta_y <- 1
alpha_z <- 1; beta_z <- 1

# Posterior parameters
alpha_y_post <- alpha_y + sum(y * n_y)
beta_y_post <- beta_y + sum((1 - y) * n_y)
alpha_z_post <- alpha_z + sum(z * n_z)
beta_z_post <- beta_z + sum((1 - z) * n_z)

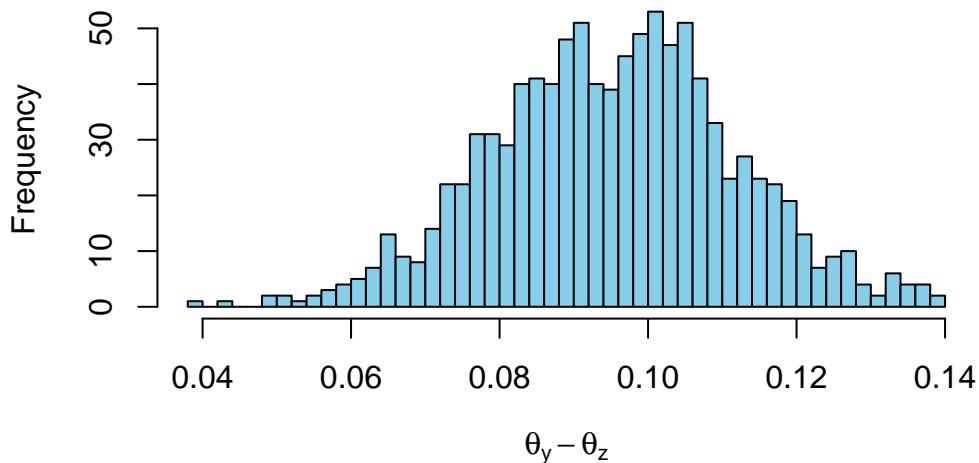
# Simulate from the posterior
set.seed(123)
theta_y_samples <- rbeta(1000, alpha_y_post, beta_y_post)
theta_z_samples <- rbeta(1000, alpha_z_post, beta_z_post)

# Compute the difference
diff_samples <- theta_y_samples - theta_z_samples

# Plot
hist(diff_samples, main = "Posterior Distribution of Difference",
     xlab = expression(theta[y] - theta[z]), breaks = 40, col = "skyblue")

```

**Posterior Distribution of Difference**



## Exercise 2

**Regression with many explanatory variables [40 points].** Table 15.2 displays data from a designed experiment for a chemical process. In using these data to illustrate various approaches to selection and estimation of regression coefficients, Marquardt and Snee (1975) assume a quadratic regression form; that is, a linear relation between the expectation of the untransformed outcome,  $y$ , and the variables  $x_1, x_2, x_3$ , their two-way interactions,  $x_1x_2, x_1x_3, x_2x_3$ , and their squares,  $x_1^2, x_2^2, x_3^2$ .

Table 1: Data from a chemical experiment, from Marquardt and Snee (1975). The first three variables are experimental manipulations, and the fourth is the outcome measurement.

Reactor temperature (°C), $x_1$	Ratio of H <sub>2</sub> to <i>n</i> -heptane (mole ratio), $x_2$	Contact time (sec), $x_3$	Conversion of <i>n</i> -heptane to acetylene (%), $y$
1300	7.5	0.0120	49.0
1300	9.0	0.0120	50.2
1300	11.0	0.0115	50.5
1300	13.5	0.0130	48.5
1300	17.0	0.0135	47.5
1300	23.0	0.0120	44.5
1200	5.3	0.0400	28.0
1200	7.5	0.0380	31.5
1200	11.0	0.0320	34.5
1200	13.5	0.0260	35.0
1200	17.0	0.0340	38.0
1200	23.0	0.0410	38.5
1100	5.3	0.0840	15.0
1100	7.5	0.0980	17.0
1100	11.0	0.0920	20.5
1100	17.0	0.0860	19.5

Data is available here: <http://www.stat.columbia.edu/~gelman/book/data/factorial.asc>

- [5 points] Fit an ordinary linear regression model (OLS frequentist model), including a constant term and the nine explanatory variables above.
- [10 points] Fit a Bayesian linear regression model with a uniform prior distribution on the constant term and a shared normal prior distribution on the coefficients of the nine variables above. If you use iterative simulation in your computations, be sure to use multiple sequences and monitor their joint convergence.
- [10 points] Discuss the differences between the inferences in (1) and (2). Interpret the differences in terms of the hierarchical variance parameter. Do you agree with Marquardt and Snee that the inferences from (1) are unacceptable?
- [8 points] Repeat (1), but with a  $t_4$  prior distribution on the nine variables.
- [7 points] Discuss other models for the regression coefficients.

## Solution

### Part 1: Ordinary Linear Regression (OLS Model)

The regression model takes the form:

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_1^2 + \beta_5 x_2^2 + \beta_6 x_3^2 + \beta_7 x_1 x_2 + \beta_8 x_2 x_3 + \beta_9 x_1 x_3 + \epsilon,$$

where  $\epsilon \sim \mathcal{N}(0, \sigma^2)$ . Using the dataset provided, I fit the model using ordinary least squares.

**R code is attached below**

The OLS regression summary provides coefficients, standard errors,  $R^2$ , and residual diagnostics.

### Part 2: Bayesian Linear Regression with Uniform Prior

I assume a uniform prior for the intercept ( $\beta_0$ ) and a shared normal prior for the other coefficients:

$$\beta_i \sim \mathcal{N}(0, \tau^2), \quad i = 1, \dots, 9,$$

with  $\tau$  as a hyperparameter to be estimated. The error variance  $\sigma^2$  is also estimated from the data.

**R code is attached below**

The posterior estimates for the coefficients  $\beta_i$ ,  $\sigma$ , and  $\tau$  are reported along with their credible intervals.

### Part 3: Comparison of Inferences

The Bayesian approach includes uncertainty in  $\tau$ , which controls the regularization of coefficients. Inferences from OLS may overfit the data if the predictors are highly correlated. The Bayesian hierarchical model mitigates this issue by shrinking the coefficients toward zero.

#### Discussion Points:

- The hierarchical variance  $\tau$  adjusts for overfitting by pooling information across predictors.
- In OLS, high multicollinearity could lead to inflated standard errors, which are not addressed without penalization.
- Marquardt and Snee highlight that OLS estimates are often unstable for such problems, making Bayesian regularization preferable.

### Part 4: OLS with $t_4$ Prior

I replace the normal prior with a Student- $t_4$  prior, which is more robust to outliers:

$$\beta_i \sim t_4(0, \sigma).$$

**R code is attached below**

The posterior estimates reflect greater robustness to deviations in the data compared to the normal prior.

## Part 5: Other Models

Alternative models include:

- Ridge regression (penalizing large coefficients).
- LASSO regression (enforcing sparsity).
- Gaussian process regression for capturing nonlinear relationships.

In this section, I explore alternative regression approaches to address the issues of multicollinearity, overfitting, and nonlinear relationships in the data.

### Ridge Regression (*Penalizing Large Coefficients*)

Ridge regression shrinks the coefficients by adding a penalty proportional to the sum of their squared values:

$$\text{Minimize: } \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p \beta_j^2.$$

**R code is attached below**

### LASSO Regression (*Enforcing Sparsity*)

LASSO regression adds a penalty proportional to the absolute value of the coefficients, encouraging sparsity:

$$\text{Minimize: } \sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^p |\beta_j|.$$

**R code is attached below**

### Gaussian Process Regression (*Capturing Nonlinear Relationships*)

Gaussian process regression (GPR) is a nonparametric method that uses a kernel function to model complex relationships between the predictors and the outcome.

**R code is attached below**

## Discussion of Models

1. **MSE (Mean Squared Error):** 71.87 indicates that the average squared difference between observed and predicted values is relatively high, suggesting the model struggles with accurately capturing the response variable.
2. **RMSE (Root Mean Squared Error):** 8.48 provides the average prediction error in the same units as the response variable, which is sizable compared to the observed values (ranging from approximately 25 to 63).
3. **MAE (Mean Absolute Error):** 6.96 shows the average absolute error is high, reflecting consistent prediction deviations from the observed values.

4.  **$R^2$  (Coefficient of Determination):** 0.26 indicates that only 26% of the variance in the observed data is explained by the model. This suggests a weak relationship between predictors and the response variable.

## Likely Issues

- The model may be underfitting the data, meaning it cannot adequately capture the underlying relationships.
- The predictors (or their transformations) may not fully explain the variability in the response.
- There may be multicollinearity or inappropriate scaling among the predictors.
- The Gaussian Process model's kernel or hyperparameters may not be optimal for this dataset.

## Suggestions to Improve the Model

### 1. Feature Engineering:

- *Standardization or Normalization:* If predictors have different scales, standardize them to mean 0 and standard deviation 1. Gaussian Processes are sensitive to input scales.
- *Interaction Terms:* Include or refine interaction terms if they are expected to influence the response.
- *Principal Component Analysis (PCA):* Reduce dimensionality to remove multicollinearity and simplify the feature space.

### 2. Hyperparameter Tuning:

Use the `kernlab::gausspr` hyperparameter options to fine-tune the kernel. For instance:

- *Test other kernels* like `polydot` (polynomial kernel) or `rbfdot` (radial basis function kernel).
- *Adjust kernel-specific parameters* like the width of the RBF kernel or degree of the polynomial.

### 3. Increase Model Flexibility:

- Use non-linear regression models, such as tree-based models (Random Forest or Gradient Boosting), which can better handle complex relationships.
- Alternatively, try a Bayesian approach to incorporate prior knowledge about the data.

### 4. Check for Outliers:

Investigate whether specific data points have undue influence on the model's performance. Remove or transform outliers as needed.

### 5. Cross-Validation:

Perform  $k$ -fold cross-validation to evaluate model stability and ensure that the performance isn't an artifact of the train-test split.

### 6. Model Comparison:

Fit other models (e.g., linear regression, ridge, lasso) to establish a baseline performance and assess whether Gaussian Processes offer a significant improvement.

Each method addresses specific limitations of OLS, providing a range of tools for regression analysis depending on the problem context.

## Conclusion

In this exercise, I explored linear regression (both frequentist and Bayesian), Ridge regression, LASSO regression, and Gaussian process regression. Bayesian methods and alternative models like Ridge and LASSO provide robust solutions to common regression issues such as multicollinearity and overfitting, while Gaussian Process regression offers flexibility for modeling nonlinear relationships.

## STATS 551 - HW 5 - Q2

```
# install.packages("rstan", repos = "https://cloud.r-project.org/")
library(rstan)

Loading required package: StanHeaders

rstan version 2.32.6 (Stan version 2.32.2)

For execution on a local, multicore CPU with excess RAM we recommend calling
options(mc.cores = parallel::detectCores()).

To avoid recompilation of unchanged Stan programs, we recommend calling
rstan_options(auto_write = TRUE)
For within-chain threading using `reduce_sum()` or `map_rect()` Stan functions,
change `threads_per_chain` option:
rstan_options(threads_per_chain = 1)

Do not specify '-march=native' in 'LOCAL_CPPFLAGS' or a Makevars file

library(mcmc)

# Load data
fact_data <- read.csv("factorial_data.csv")

# Prepare data
fact_data$temperature_sq <- fact_data$temperature^2
fact_data$ratio_sq <- fact_data$ratio^2
fact_data$contact_sq <- fact_data$contact^2
fact_data$temperature_ratio <- fact_data$temperature * fact_data$ratio
fact_data$ratio_contact <- fact_data$ratio * fact_data$contact
fact_data$temperature_contact <- fact_data$temperature * fact_data$contact

# Fit OLS model
ols_model <- lm(conversion ~ temperature + ratio + contact +
                  temperature_sq + ratio_sq + contact_sq +
                  temperature_ratio + ratio_contact + temperature_contact,
                  data = fact_data)
summary(ols_model)
```

Call:

```
lm(formula = conversion ~ temperature + ratio + contact + temperature_sq +
```

```

ratio_sq + contact_sq + temperature_ratio + ratio_contact +
temperature_contact, data = fact_data)

```

Residuals:

Min	1Q	Median	3Q	Max
-1.4824	-0.5007	0.1549	0.3887	1.3776

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-7.849e+03	4.403e+03	-1.783	0.12493
temperature	1.178e+01	6.851e+00	1.719	0.13638
ratio	3.164e+01	6.042e+00	5.237	0.00194 **
contact	2.634e+04	1.467e+04	1.795	0.12277
temperature_sq	-4.390e-03	2.662e-03	-1.649	0.15021
ratio_sq	-4.352e-02	1.641e-02	-2.652	0.03791 *
contact_sq	-2.087e+04	1.081e+04	-1.931	0.10173
temperature_ratio	-2.299e-02	4.509e-03	-5.098	0.00223 **
ratio_contact	-5.768e+01	1.297e+01	-4.445	0.00435 **
temperature_contact	-2.004e+01	1.157e+01	-1.732	0.13398
---				
Signif. codes:	0 ***	0.001 **	0.01 *	0.05 .
	'	'	'	'
	'	'	'	'

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.266 on 6 degrees of freedom

Multiple R-squared: 0.9959, Adjusted R-squared: 0.9898

F-statistic: 162.3 on 9 and 6 DF, p-value: 1.816e-06

Bayesian Model has been specified in STAN file

```

# Prepare data for Stan
X <- scale(fact_data[, c("temperature", "ratio", "contact",
                         "temperature_sq", "ratio_sq", "contact_sq",
                         "temperature_ratio", "ratio_contact", "temperature_contact")])

y <- fact_data$conversion

stan_data <- list(
  N = nrow(fact_data),
  K = ncol(X),
  X = X,
  y = y
)

# Fit Bayesian model
fit <- stan(file = "Q2_Bayesian_Model.stan", data = stan_data, iter = 8000, chains = 4,
            control = list(adapt_delta = 0.99, max_treedepth = 15))

```

SAMPLING FOR MODEL 'anon\_model' NOW (CHAIN 1).

Chain 1:

```
Chain 1: Gradient evaluation took 7.9e-05 seconds
Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.79 seconds.
Chain 1: Adjust your expectations accordingly!
Chain 1:
Chain 1:
Chain 1: Iteration:    1 / 8000 [  0%] (Warmup)
Chain 1: Iteration:  800 / 8000 [ 10%] (Warmup)
Chain 1: Iteration: 1600 / 8000 [ 20%] (Warmup)
Chain 1: Iteration: 2400 / 8000 [ 30%] (Warmup)
Chain 1: Iteration: 3200 / 8000 [ 40%] (Warmup)
Chain 1: Iteration: 4000 / 8000 [ 50%] (Warmup)
Chain 1: Iteration: 4001 / 8000 [ 50%] (Sampling)
Chain 1: Iteration: 4800 / 8000 [ 60%] (Sampling)
Chain 1: Iteration: 5600 / 8000 [ 70%] (Sampling)
Chain 1: Iteration: 6400 / 8000 [ 80%] (Sampling)
Chain 1: Iteration: 7200 / 8000 [ 90%] (Sampling)
Chain 1: Iteration: 8000 / 8000 [100%] (Sampling)
Chain 1:
Chain 1: Elapsed Time: 3.619 seconds (Warm-up)
Chain 1:           3.058 seconds (Sampling)
Chain 1:           6.677 seconds (Total)
Chain 1:
```

```
SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
Chain 2:
Chain 2: Gradient evaluation took 6.7e-05 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.67 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration:    1 / 8000 [  0%] (Warmup)
Chain 2: Iteration:  800 / 8000 [ 10%] (Warmup)
Chain 2: Iteration: 1600 / 8000 [ 20%] (Warmup)
Chain 2: Iteration: 2400 / 8000 [ 30%] (Warmup)
Chain 2: Iteration: 3200 / 8000 [ 40%] (Warmup)
Chain 2: Iteration: 4000 / 8000 [ 50%] (Warmup)
Chain 2: Iteration: 4001 / 8000 [ 50%] (Sampling)
Chain 2: Iteration: 4800 / 8000 [ 60%] (Sampling)
Chain 2: Iteration: 5600 / 8000 [ 70%] (Sampling)
Chain 2: Iteration: 6400 / 8000 [ 80%] (Sampling)
Chain 2: Iteration: 7200 / 8000 [ 90%] (Sampling)
Chain 2: Iteration: 8000 / 8000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 3.563 seconds (Warm-up)
Chain 2:           4.865 seconds (Sampling)
Chain 2:           8.428 seconds (Total)
Chain 2:
```

```
SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 3).  
Chain 3:  
Chain 3: Gradient evaluation took 1.8e-05 seconds  
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.18 seconds.  
Chain 3: Adjust your expectations accordingly!  
Chain 3:  
Chain 3:  
Chain 3: Iteration: 1 / 8000 [  0%] (Warmup)  
Chain 3: Iteration: 800 / 8000 [ 10%] (Warmup)  
Chain 3: Iteration: 1600 / 8000 [ 20%] (Warmup)  
Chain 3: Iteration: 2400 / 8000 [ 30%] (Warmup)  
Chain 3: Iteration: 3200 / 8000 [ 40%] (Warmup)  
Chain 3: Iteration: 4000 / 8000 [ 50%] (Warmup)  
Chain 3: Iteration: 4001 / 8000 [ 50%] (Sampling)  
Chain 3: Iteration: 4800 / 8000 [ 60%] (Sampling)  
Chain 3: Iteration: 5600 / 8000 [ 70%] (Sampling)  
Chain 3: Iteration: 6400 / 8000 [ 80%] (Sampling)  
Chain 3: Iteration: 7200 / 8000 [ 90%] (Sampling)  
Chain 3: Iteration: 8000 / 8000 [100%] (Sampling)  
Chain 3:  
Chain 3: Elapsed Time: 2.634 seconds (Warm-up)  
Chain 3: 2.381 seconds (Sampling)  
Chain 3: 5.015 seconds (Total)  
Chain 3:
```

```
SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 4).  
Chain 4:  
Chain 4: Gradient evaluation took 3.5e-05 seconds  
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.35 seconds.  
Chain 4: Adjust your expectations accordingly!  
Chain 4:  
Chain 4:  
Chain 4: Iteration: 1 / 8000 [  0%] (Warmup)  
Chain 4: Iteration: 800 / 8000 [ 10%] (Warmup)  
Chain 4: Iteration: 1600 / 8000 [ 20%] (Warmup)  
Chain 4: Iteration: 2400 / 8000 [ 30%] (Warmup)  
Chain 4: Iteration: 3200 / 8000 [ 40%] (Warmup)  
Chain 4: Iteration: 4000 / 8000 [ 50%] (Warmup)  
Chain 4: Iteration: 4001 / 8000 [ 50%] (Sampling)  
Chain 4: Iteration: 4800 / 8000 [ 60%] (Sampling)  
Chain 4: Iteration: 5600 / 8000 [ 70%] (Sampling)  
Chain 4: Iteration: 6400 / 8000 [ 80%] (Sampling)  
Chain 4: Iteration: 7200 / 8000 [ 90%] (Sampling)  
Chain 4: Iteration: 8000 / 8000 [100%] (Sampling)  
Chain 4:  
Chain 4: Elapsed Time: 2.892 seconds (Warm-up)
```

```

Chain 4:           3.179 seconds (Sampling)
Chain 4:           6.071 seconds (Total)
Chain 4:

print(fit)

Inference for Stan model: anon_model.
4 chains, each with iter=8000; warmup=4000; thin=1;
post-warmup draws per chain=4000, total post-warmup draws=16000.

          mean se_mean    sd  2.5%   25%   50%   75% 97.5% n_eff Rhat
beta_0    34.73    0.01 0.89  32.73  34.30  34.83  35.29  36.14  6451    1
beta[1]    4.86    0.04 3.81  -2.61   2.32   4.89   7.42  12.29 11705    1
beta[2]    1.73    0.04 3.90  -5.98  -0.82   1.75   4.35   9.37 11114    1
beta[3]   -2.38    0.04 4.26 -10.83  -5.23  -2.37   0.48   5.89 12565    1
beta[4]    4.85    0.03 3.64  -2.23   2.39   4.87   7.32  11.98 11429    1
beta[5]   -2.62    0.03 2.96  -8.32  -4.63  -2.67  -0.69   3.35 12234    1
beta[6]    0.15    0.03 3.04  -5.88  -1.91   0.18   2.22   5.96 12204    1
beta[7]    0.55    0.04 3.87  -6.99  -2.09   0.55   3.17   8.16 11404    1
beta[8]    2.67    0.02 1.89  -1.15   1.49   2.73   3.91   6.22 11650    1
beta[9]   -2.27    0.04 4.15 -10.42  -5.06  -2.28   0.55   5.92 13035    1
sigma     2.83    0.01 0.85   1.79   2.29   2.67   3.17   4.88  4629    1
lp__   -51.90    0.04 2.85 -58.71 -53.54 -51.47 -49.83 -47.60  4075    1

```

Samples were drawn using NUTS(diag\_e) at Mon Dec 9 19:37:17 2024.  
For each parameter, n\_eff is a crude measure of effective sample size,  
and Rhat is the potential scale reduction factor on split chains (at  
convergence, Rhat=1).

```

library(mcmc)
library(MCMCpack)

```

Loading required package: coda

Attaching package: 'coda'

The following object is masked from 'package:rstan':

traceplot

Loading required package: MASS

##

## Markov Chain Monte Carlo Package (MCMCpack)

## Copyright (C) 2003-2024 Andrew D. Martin, Kevin M. Quinn, and Jong Hee Park

##

## Support provided by the U.S. National Science Foundation

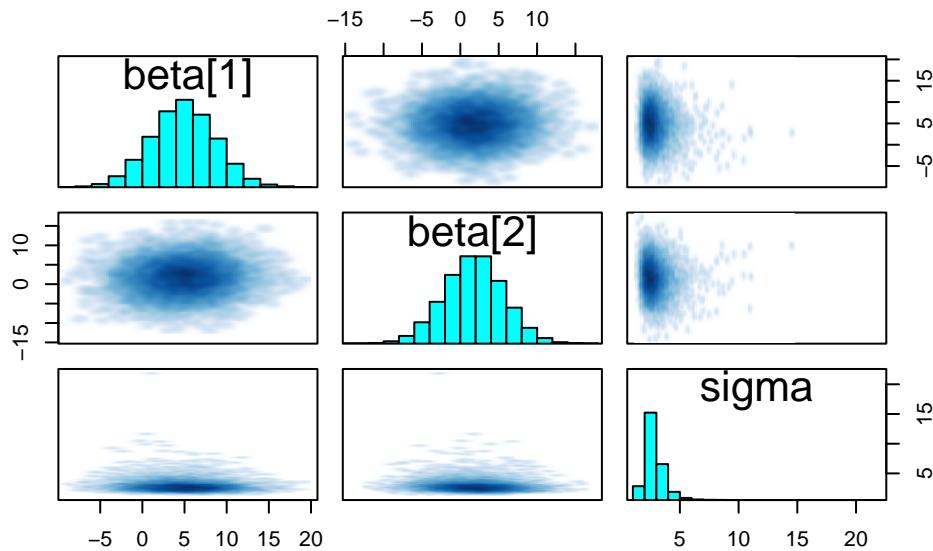
## (Grants SES-0350646 and SES-0350613)

##

```
library(bayesplot)
```

```
This is bayesplot version 1.11.1
- Online documentation and vignettes at mc-stan.org/bayesplot
- bayesplot theme set to bayesplot::theme_default()
  * Does _not_ affect other ggplot2 plots
  * See ?bayesplot_theme_set for details on theme setting
pairs(fit, pars = c("beta[1]", "beta[2]", "sigma"))
```

```
Warning in par(usr): argument 1 does not name a graphical parameter
Warning in par(usr): argument 1 does not name a graphical parameter
Warning in par(usr): argument 1 does not name a graphical parameter
```



```
summary(fit) # R-hat and ESS for all parameters
```

```
$summary
```

	mean	se_mean	sd	2.5%	25%	50%
beta_0	34.7297463	0.01113756	0.8945624	32.733778	34.298090	34.8262293
beta[1]	4.8552738	0.03519004	3.8072753	-2.614091	2.316751	4.8857533
beta[2]	1.7320465	0.03694920	3.8953491	-5.984051	-0.823760	1.7461111
beta[3]	-2.3777530	0.03797206	4.2563919	-10.831783	-5.227896	-2.3695265
beta[4]	4.8518944	0.03408177	3.6435397	-2.232770	2.391829	4.8655362
beta[5]	-2.6206258	0.02678754	2.9629367	-8.323180	-4.628194	-2.6665177
beta[6]	0.1519888	0.02752130	3.0403558	-5.883436	-1.910409	0.1839381
beta[7]	0.5549907	0.03624296	3.8703860	-6.988192	-2.092964	0.5475916
beta[8]	2.6700401	0.01750386	1.8892816	-1.146375	1.486675	2.7294960
beta[9]	-2.2704435	0.03638537	4.1541972	-10.422348	-5.056305	-2.2799762
sigma	2.8333848	0.01243438	0.8459848	1.785776	2.285217	2.6736114
lp__	-51.9030069	0.04464894	2.8502582	-58.714101	-53.535719	-51.4694965
	75%	97.5%	n_eff	Rhat		
beta_0	35.2924994	36.138970	6451.210	1.0004992		

```

beta[1]    7.4243648 12.288793 11705.476 1.0001130
beta[2]    4.3489012 9.367303 11114.312 1.0000048
beta[3]    0.4844779 5.889531 12564.783 0.9998654
beta[4]    7.3192377 11.979278 11428.856 0.9998839
beta[5]   -0.6927525 3.354567 12234.296 1.0000354
beta[6]    2.2152365 5.956759 12204.248 0.9998701
beta[7]    3.1713646 8.160227 11404.106 1.0001006
beta[8]    3.9078386 6.215903 11649.991 1.0003749
beta[9]    0.5536505 5.916787 13035.290 1.0000931
sigma     3.1677524 4.882980 4628.891 1.0006960
lp__     -49.8277653 -47.603611 4075.174 1.0002706

```

```

$c_summary
, , chains = chain:1

```

stats						
parameter	mean	sd	2.5%	25%	50%	75%
beta_0	34.7177249	0.9365730	32.702377	34.2938962	34.8085484	35.2757986
beta[1]	4.7406714	3.7963801	-2.670617	2.2083311	4.7666850	7.3137455
beta[2]	1.6522446	3.8092930	-5.611975	-0.9374642	1.6822550	4.2701698
beta[3]	-2.3234761	4.1755249	-10.673906	-5.1028587	-2.2990161	0.5205173
beta[4]	4.8798796	3.6154727	-2.057805	2.4397782	4.9103745	7.3358339
beta[5]	-2.5587450	2.9830076	-8.457626	-4.5270235	-2.5039310	-0.6061142
beta[6]	0.1760979	3.0325927	-5.768303	-1.8936758	0.1758535	2.1910699
beta[7]	0.6306519	3.8394264	-6.911933	-2.0160860	0.6948570	3.2666335
beta[8]	2.6033954	1.9422002	-1.245412	1.4271616	2.6787079	3.9022772
beta[9]	-2.3551068	4.1258642	-10.266973	-5.1955233	-2.3389446	0.4660599
sigma	2.8597624	0.9194119	1.776980	2.2866076	2.6906545	3.2032078
lp__	-51.9405860	2.9244822	-58.906566	-53.5151203	-51.4776876	-49.8405467

stats	
parameter	97.5%
beta_0	36.142194
beta[1]	12.044407
beta[2]	8.997473
beta[3]	5.775922
beta[4]	11.973607
beta[5]	3.299188
beta[6]	6.165061
beta[7]	8.004673
beta[8]	6.171223
beta[9]	5.674762
sigma	4.989782
lp__	-47.657699

```

, , chains = chain:2

```

```

stats

```

```

parameter      mean       sd     2.5%     25%     50%     75%
beta_0    34.7614768 0.8229985 32.805474 34.3058273 34.83548726 35.3134976
beta[1]   4.9053105 3.8421347 -2.619201 2.3503256 4.96963173 7.4713004
beta[2]   1.7065561 3.9245730 -6.081358 -0.8470594 1.73467015 4.3245912
beta[3]   -2.4105081 4.3378190 -10.832358 -5.3875597 -2.40427162 0.5495958
beta[4]   4.8375196 3.6541019 -2.172926 2.3287314 4.84704099 7.3548487
beta[5]   -2.6488415 2.9412308 -8.220191 -4.5974979 -2.73675150 -0.7451877
beta[6]   0.1172412 3.0409822 -5.775193 -1.9500532 0.07426162 2.1640043
beta[7]   0.5917296 3.8026574 -6.854740 -1.9905874 0.58829164 3.1452592
beta[8]   2.7305990 1.8402158 -1.064466 1.5941909 2.79594571 3.9510504
beta[9]   -2.2198983 4.1650103 -10.199662 -5.0359263 -2.24236500 0.5887050
sigma    2.8045110 0.7675432 1.776150 2.2761590 2.66279679 3.1582041
lp__    -51.8377766 2.7445212 -58.354917 -53.4582467 -51.45553924 -49.8254728

stats
parameter      97.5%
beta_0    36.140538
beta[1]   12.320796
beta[2]   9.303106
beta[3]   6.112723
beta[4]   11.920293
beta[5]   3.223650
beta[6]   6.093142
beta[7]   8.018332
beta[8]   6.090996
beta[9]   6.068555
sigma    4.780528
lp__    -47.570798

, , chains = chain:3

stats
parameter      mean       sd     2.5%     25%     50%     75%
beta_0    34.7236057 0.9323210 32.735528 34.3017643 34.8291550 35.3071243
beta[1]   4.8935535 3.7677338 -2.513049 2.4196115 4.8684860 7.5307174
beta[2]   1.8379567 3.9411011 -5.990415 -0.7359934 1.9154666 4.4501644
beta[3]   -2.3443053 4.2574448 -10.809536 -5.1548448 -2.3201472 0.4335684
beta[4]   4.8314493 3.6318472 -2.401941 2.3984712 4.8209202 7.2510224
beta[5]   -2.6763404 2.9737784 -8.218261 -4.7151679 -2.7807579 -0.7298958
beta[6]   0.1889162 3.0437657 -5.961031 -1.8559579 0.2718271 2.2985085
beta[7]   0.4797046 3.8816032 -7.011593 -2.1650024 0.4415692 3.0843970
beta[8]   2.6704744 1.9025237 -1.151504 1.4442976 2.7412646 3.8810180
beta[9]   -2.3270318 4.1796215 -10.550804 -5.1480724 -2.3642692 0.5273778
sigma    2.8493492 0.8712253 1.783957 2.2865588 2.6776911 3.1727337
lp__    -51.9619543 2.8594758 -58.811998 -53.6156110 -51.5572410 -49.8809371

stats
parameter      97.5%
beta_0    36.126706

```

```

beta[1] 12.133217
beta[2] 9.698081
beta[3] 5.969201
beta[4] 11.974302
beta[5] 3.427575
beta[6] 5.820222
beta[7] 8.190510
beta[8] 6.329762
beta[9] 5.920550
sigma    4.882107
lp__    -47.641626

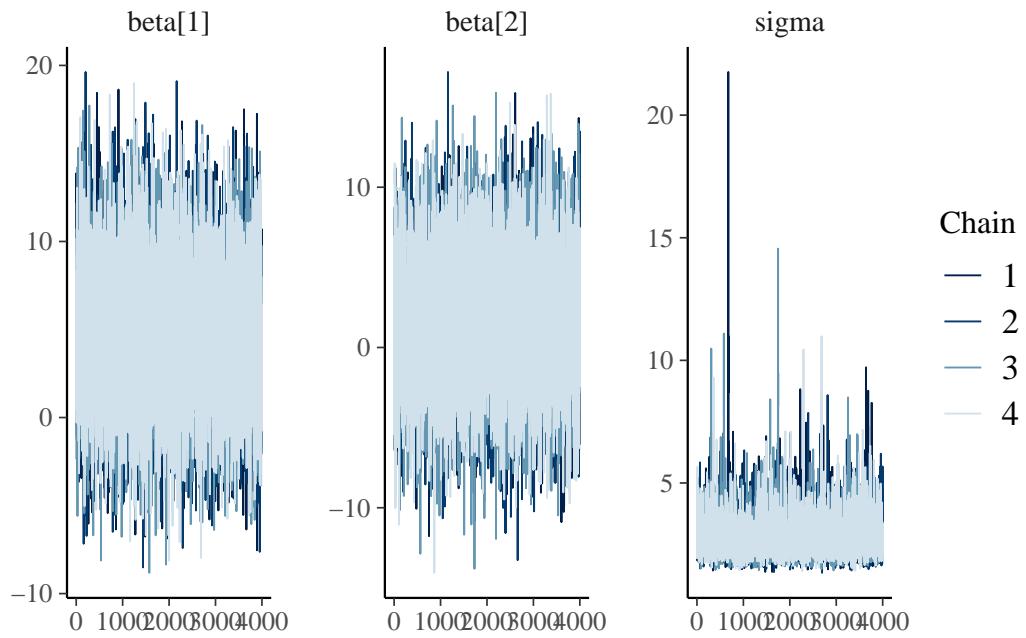
, , chains = chain:4

      stats
parameter   mean      sd     2.5%    25%    50%    75%
beta_0    34.7161778 0.8811792 32.707008 34.287996 34.8262603 35.2677555
beta[1]   4.8815599 3.8215360 -2.587008  2.310759  4.9004122 7.3869533
beta[2]   1.7314285 3.9042101 -6.030005 -0.775548  1.6946440 4.3364476
beta[3]  -2.4327225 4.2538687 -11.117932 -5.281092 -2.4588035 0.4671506
beta[4]   4.8587293 3.6736377 -2.383022  2.389039  4.8704640 7.3251743
beta[5]  -2.5985763 2.9532714 -8.235493 -4.615309 -2.6480486 -0.6809177
beta[6]   0.1256996 3.0445760 -6.065927 -1.899410  0.1822548 2.1854205
beta[7]   0.5178768 3.9558089 -7.152622 -2.190244  0.4970693 3.1806743
beta[8]   2.6756914 1.8692089 -1.098922  1.509874  2.7091948 3.8909683
beta[9]  -2.1797369 4.1451059 -10.572708 -4.869264 -2.1916559 0.6146421
sigma    2.8199168 0.8172256  1.805992  2.290822  2.6642732 3.1297125
lp__    -51.8717108 2.8688435 -58.791867 -53.530061 -51.3884984 -49.7645155

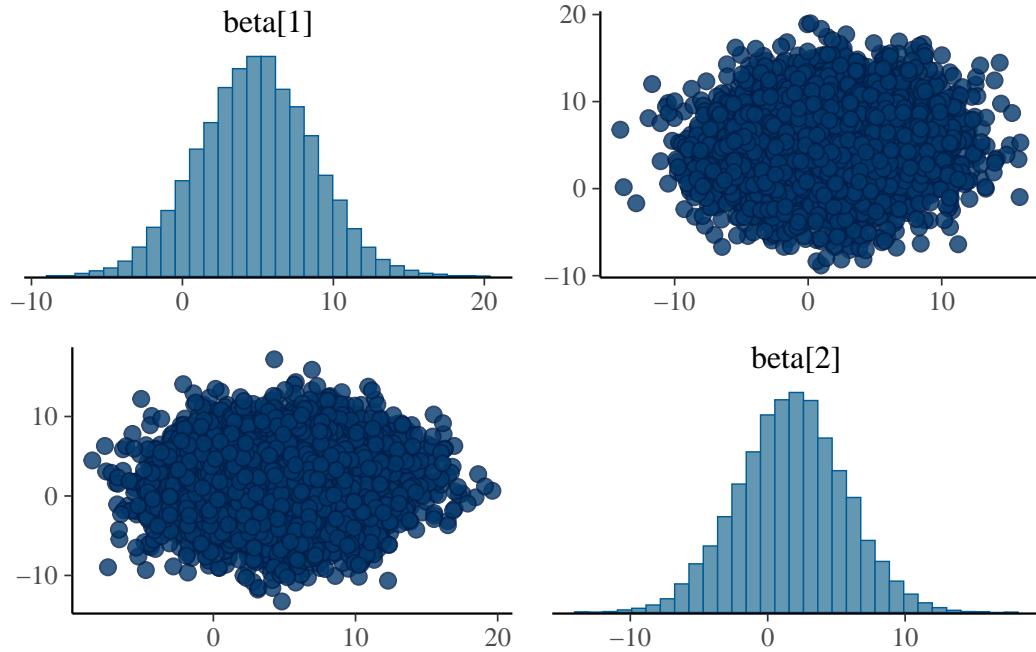
      stats
parameter   97.5%
beta_0    36.111530
beta[1]   12.741900
beta[2]   9.443303
beta[3]   5.797287
beta[4]  12.046004
beta[5]   3.471923
beta[6]   5.854435
beta[7]   8.355726
beta[8]   6.306088
beta[9]   5.948398
sigma    4.848801
lp__    -47.582895

posterior <- as.array(fit)
mcmc_trace(posterior, pars = c("beta[1]", "beta[2]", "sigma")) # Trace plots

```



```
mcmc_pairs(posterior, pars = c("beta[1]", "beta[2]")) # Joint posterior diagnostics
```



```
fit_t4 <- stan(file = "Q2-Bayesian-Model-t4-prior.stan",
                 data = stan_data,
                 iter = 8000,
                 chains = 4,
                 control = list(adapt_delta = 0.9999, max_treedepth = 15))
```

SAMPLING FOR MODEL 'anon\_model' NOW (CHAIN 1).  
 Chain 1:  
 Chain 1: Gradient evaluation took 8.2e-05 seconds

```
Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.82 seconds.  
Chain 1: Adjust your expectations accordingly!  
Chain 1:  
Chain 1:  
Chain 1: Iteration: 1 / 8000 [ 0%] (Warmup)  
Chain 1: Iteration: 800 / 8000 [ 10%] (Warmup)  
Chain 1: Iteration: 1600 / 8000 [ 20%] (Warmup)  
Chain 1: Iteration: 2400 / 8000 [ 30%] (Warmup)  
Chain 1: Iteration: 3200 / 8000 [ 40%] (Warmup)  
Chain 1: Iteration: 4000 / 8000 [ 50%] (Warmup)  
Chain 1: Iteration: 4001 / 8000 [ 50%] (Sampling)  
Chain 1: Iteration: 4800 / 8000 [ 60%] (Sampling)  
Chain 1: Iteration: 5600 / 8000 [ 70%] (Sampling)  
Chain 1: Iteration: 6400 / 8000 [ 80%] (Sampling)  
Chain 1: Iteration: 7200 / 8000 [ 90%] (Sampling)  
Chain 1: Iteration: 8000 / 8000 [100%] (Sampling)  
Chain 1:  
Chain 1: Elapsed Time: 21.216 seconds (Warm-up)  
Chain 1: 6.185 seconds (Sampling)  
Chain 1: 27.401 seconds (Total)  
Chain 1:
```

```
SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).  
Chain 2:  
Chain 2: Gradient evaluation took 1.2e-05 seconds  
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.12 seconds.  
Chain 2: Adjust your expectations accordingly!  
Chain 2:  
Chain 2:  
Chain 2: Iteration: 1 / 8000 [ 0%] (Warmup)  
Chain 2: Iteration: 800 / 8000 [ 10%] (Warmup)  
Chain 2: Iteration: 1600 / 8000 [ 20%] (Warmup)  
Chain 2: Iteration: 2400 / 8000 [ 30%] (Warmup)  
Chain 2: Iteration: 3200 / 8000 [ 40%] (Warmup)  
Chain 2: Iteration: 4000 / 8000 [ 50%] (Warmup)  
Chain 2: Iteration: 4001 / 8000 [ 50%] (Sampling)  
Chain 2: Iteration: 4800 / 8000 [ 60%] (Sampling)  
Chain 2: Iteration: 5600 / 8000 [ 70%] (Sampling)  
Chain 2: Iteration: 6400 / 8000 [ 80%] (Sampling)  
Chain 2: Iteration: 7200 / 8000 [ 90%] (Sampling)  
Chain 2: Iteration: 8000 / 8000 [100%] (Sampling)  
Chain 2:  
Chain 2: Elapsed Time: 10.958 seconds (Warm-up)  
Chain 2: 1.292 seconds (Sampling)  
Chain 2: 12.25 seconds (Total)  
Chain 2:
```

```
SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 3).  
Chain 3:  
Chain 3: Gradient evaluation took 1e-05 seconds  
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.1 seconds.  
Chain 3: Adjust your expectations accordingly!  
Chain 3:  
Chain 3:  
Chain 3: Iteration: 1 / 8000 [ 0%] (Warmup)  
Chain 3: Iteration: 800 / 8000 [ 10%] (Warmup)  
Chain 3: Iteration: 1600 / 8000 [ 20%] (Warmup)  
Chain 3: Iteration: 2400 / 8000 [ 30%] (Warmup)  
Chain 3: Iteration: 3200 / 8000 [ 40%] (Warmup)  
Chain 3: Iteration: 4000 / 8000 [ 50%] (Warmup)  
Chain 3: Iteration: 4001 / 8000 [ 50%] (Sampling)  
Chain 3: Iteration: 4800 / 8000 [ 60%] (Sampling)  
Chain 3: Iteration: 5600 / 8000 [ 70%] (Sampling)  
Chain 3: Iteration: 6400 / 8000 [ 80%] (Sampling)  
Chain 3: Iteration: 7200 / 8000 [ 90%] (Sampling)  
Chain 3: Iteration: 8000 / 8000 [100%] (Sampling)  
Chain 3:  
Chain 3: Elapsed Time: 17.072 seconds (Warm-up)  
Chain 3: 2.602 seconds (Sampling)  
Chain 3: 19.674 seconds (Total)  
Chain 3:
```

```
SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 4).  
Chain 4:  
Chain 4: Gradient evaluation took 9.4e-05 seconds  
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.94 seconds.  
Chain 4: Adjust your expectations accordingly!  
Chain 4:  
Chain 4:  
Chain 4: Iteration: 1 / 8000 [ 0%] (Warmup)  
Chain 4: Iteration: 800 / 8000 [ 10%] (Warmup)  
Chain 4: Iteration: 1600 / 8000 [ 20%] (Warmup)  
Chain 4: Iteration: 2400 / 8000 [ 30%] (Warmup)  
Chain 4: Iteration: 3200 / 8000 [ 40%] (Warmup)  
Chain 4: Iteration: 4000 / 8000 [ 50%] (Warmup)  
Chain 4: Iteration: 4001 / 8000 [ 50%] (Sampling)  
Chain 4: Iteration: 4800 / 8000 [ 60%] (Sampling)  
Chain 4: Iteration: 5600 / 8000 [ 70%] (Sampling)  
Chain 4: Iteration: 6400 / 8000 [ 80%] (Sampling)  
Chain 4: Iteration: 7200 / 8000 [ 90%] (Sampling)  
Chain 4: Iteration: 8000 / 8000 [100%] (Sampling)  
Chain 4:  
Chain 4: Elapsed Time: 10.608 seconds (Warm-up)  
Chain 4: 3.074 seconds (Sampling)
```

```

Chain 4:           13.682 seconds (Total)
Chain 4:

Warning: There were 14403 divergent transitions after warmup. See
https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
to find out why this is a problem and how to eliminate them.

Warning: Examine the pairs() plot to diagnose sampling problems

```

```
print(fit_t4)
```

```

Inference for Stan model: anon_model.
4 chains, each with iter=8000; warmup=4000; thin=1;
post-warmup draws per chain=4000, total post-warmup draws=16000.

```

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
beta_0	7.71	0.07	2.35	1.33	6.84	8.43	9.37	9.95	1050	1.00
beta_raw[1]	0.36	0.07	1.67	-2.45	-0.50	0.24	1.04	4.10	534	1.01
beta_raw[2]	0.11	0.05	1.39	-2.78	-0.64	0.10	0.87	2.96	878	1.00
beta_raw[3]	-0.23	0.04	1.25	-2.97	-0.92	-0.15	0.53	2.12	1126	1.00
beta_raw[4]	0.34	0.06	1.49	-2.16	-0.48	0.23	1.02	3.51	670	1.01
beta_raw[5]	0.09	0.04	1.31	-2.53	-0.68	0.09	0.82	2.86	1190	1.00
beta_raw[6]	-0.28	0.04	1.36	-3.26	-0.97	-0.22	0.50	2.32	1001	1.00
beta_raw[7]	0.06	0.04	1.35	-2.76	-0.66	0.08	0.81	2.66	1085	1.00
beta_raw[8]	-0.24	0.05	1.48	-3.16	-0.92	-0.16	0.60	2.34	757	1.01
beta_raw[9]	-0.33	0.05	1.40	-3.75	-1.05	-0.24	0.49	2.19	761	1.00
sigma	32.46	0.20	6.85	22.41	27.48	31.25	36.14	48.67	1223	1.00
beta[1]	0.36	0.07	1.67	-2.45	-0.50	0.24	1.04	4.10	534	1.01
beta[2]	0.11	0.05	1.39	-2.78	-0.64	0.10	0.87	2.96	878	1.00
beta[3]	-0.23	0.04	1.25	-2.97	-0.92	-0.15	0.53	2.12	1126	1.00
beta[4]	0.34	0.06	1.49	-2.16	-0.48	0.23	1.02	3.51	670	1.01
beta[5]	0.09	0.04	1.31	-2.53	-0.68	0.09	0.82	2.86	1190	1.00
beta[6]	-0.28	0.04	1.36	-3.26	-0.97	-0.22	0.50	2.32	1001	1.00
beta[7]	0.06	0.04	1.35	-2.76	-0.66	0.08	0.81	2.66	1085	1.00
beta[8]	-0.24	0.05	1.48	-3.16	-0.92	-0.16	0.60	2.34	757	1.01
beta[9]	-0.33	0.05	1.40	-3.75	-1.05	-0.24	0.49	2.19	761	1.00
lp__	-65.90	0.10	3.16	-73.09	-67.76	-65.49	-63.56	-60.98	925	1.01

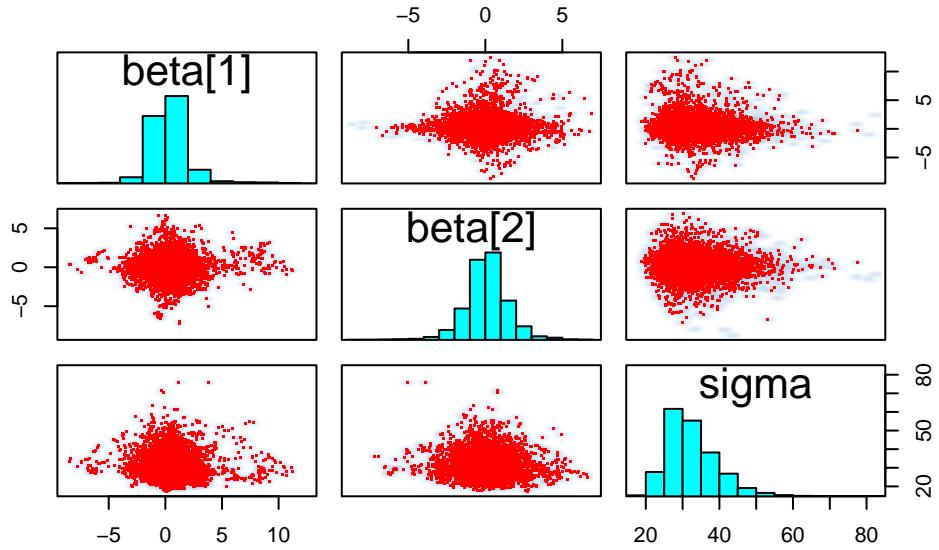
Samples were drawn using NUTS(diag\_e) at Mon Dec 9 19:41:21 2024.  
For each parameter, n\_eff is a crude measure of effective sample size,  
and Rhat is the potential scale reduction factor on split chains (at  
convergence, Rhat=1).

```
pairs(fit_t4, pars = c("beta[1]", "beta[2]", "sigma"))
```

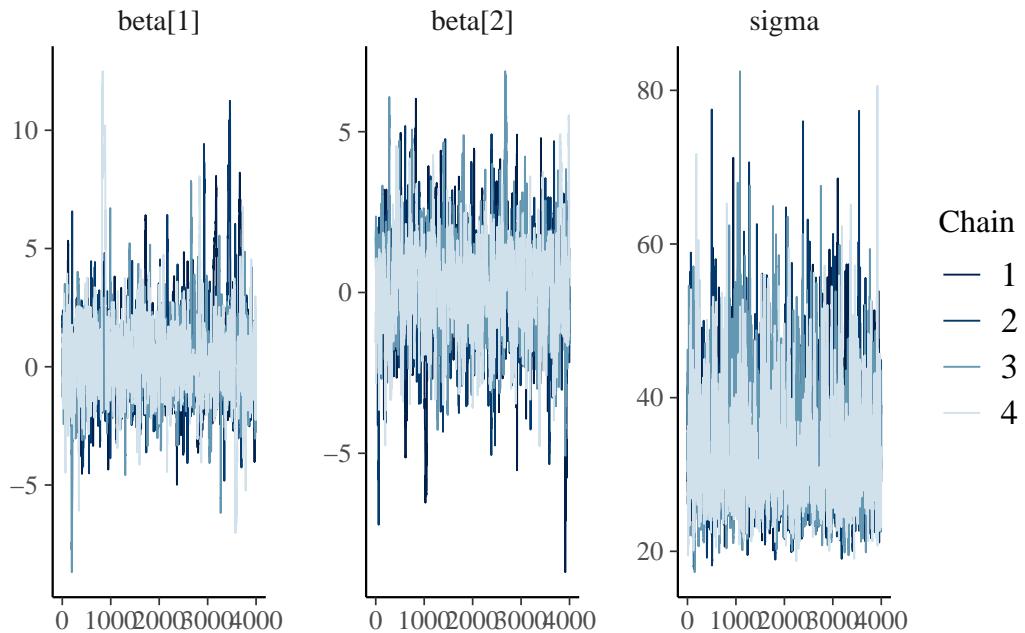
```

Warning in par(usr): argument 1 does not name a graphical parameter
Warning in par(usr): argument 1 does not name a graphical parameter
Warning in par(usr): argument 1 does not name a graphical parameter

```



```
posterior <- as.array(fit_t4)
mcmc_trace(posterior, pars = c("beta[1]", "beta[2]", "sigma"))
```



```
library(glmnet)
```

Loading required package: Matrix

Loaded glmnet 4.1-8

```
# Simulated example dataset
set.seed(123)
n <- 100
data <- data.frame(
  ratio = rnorm(n, 1300, 100),
  contact = runif(n, 5, 23),
  conversion = rnorm(n, 0.03, 0.01),
```

```

y = rnorm(n, 40, 10)
)

# Add interaction and squared terms
data$ratio_sq <- data$ratio^2
data$contact_sq <- data$contact^2
data$conversion_sq <- data$conversion^2
data$ratiocontact <- data$ratio * data$contact
data$contactconversion <- data$contact * data$conversion
data$ratioconversion <- data$ratio * data$conversion

# Prepare matrix for glmnet
X <- as.matrix(data[, c("ratio", "contact", "conversion", "ratio_sq", "contact_sq",
                       "conversion_sq", "ratiocontact", "contactconversion",
                       "ratioconversion")])
y <- data$y

# Fit ridge regression
ridge_model <- glmnet(X, y, alpha = 0) # alpha = 0 for ridge regression
ridge_cv <- cv.glmnet(X, y, alpha = 0) # Cross-validation to find best lambda
ridge_best_lambda <- ridge_cv$lambda.min

# Extract coefficients at best lambda
ridge_coefficients <- coef(ridge_cv, s = ridge_best_lambda)
print(ridge_coefficients)

```

```

10 x 1 sparse Matrix of class "dgCMatrix"
           s1
(Intercept) 4.125752e+01
ratio        5.145622e-39
contact      -3.678847e-38
conversion   5.997559e-35
ratio_sq     1.828757e-42
contact_sq   -4.366979e-40
conversion_sq 6.464200e-34
ratiocontact -3.078151e-42
contactconversion 9.398377e-37
ratioconversion 5.034607e-38

```

```

# install.packages("brms")
library(brms)

```

Loading required package: Rcpp

Loading 'brms' package (version 2.22.0). Useful instructions  
can be found by typing `help('brms')`. A more detailed introduction  
to the package is available through `vignette('brms_overview')`.

```

Attaching package: 'brms'

The following object is masked from 'package:bayesplot':

  rhat

The following objects are masked from 'package:MCMCpack':

  ddirichlet, rdirichlet

The following object is masked from 'package:rstan':

  loo

The following object is masked from 'package:stats':

  ar

# Fitting a Bayesian regression model
posterior_model <- brm(
  y ~ ratio + contact + conversion + ratio_sq + contact_sq + conversion_sq +
    ratiocontact + contactconversion + ratioconversion,
  data = data,
  family = gaussian()
)

```

Compiling Stan program...

Start sampling

SAMPLING FOR MODEL 'anon\_model' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 5.9e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.59 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [ 0%] (Warmup)

Chain 1: Iteration: 200 / 2000 [ 10%] (Warmup)

Chain 1: Iteration: 400 / 2000 [ 20%] (Warmup)

Chain 1: Iteration: 600 / 2000 [ 30%] (Warmup)

Chain 1: Iteration: 800 / 2000 [ 40%] (Warmup)

Chain 1: Iteration: 1000 / 2000 [ 50%] (Warmup)

Chain 1: Iteration: 1001 / 2000 [ 50%] (Sampling)

Chain 1: Iteration: 1200 / 2000 [ 60%] (Sampling)

Chain 1: Iteration: 1400 / 2000 [ 70%] (Sampling)

Chain 1: Iteration: 1600 / 2000 [ 80%] (Sampling)

Chain 1: Iteration: 1800 / 2000 [ 90%] (Sampling)

Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 3.279 seconds (Warm-up)

```

Chain 1:           1.005 seconds (Sampling)
Chain 1:           4.284 seconds (Total)
Chain 1:

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
Chain 2:
Chain 2: Gradient evaluation took 9e-06 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.09 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration: 1 / 2000 [  0%] (Warmup)
Chain 2: Iteration: 200 / 2000 [ 10%] (Warmup)
Chain 2: Iteration: 400 / 2000 [ 20%] (Warmup)
Chain 2: Iteration: 600 / 2000 [ 30%] (Warmup)
Chain 2: Iteration: 800 / 2000 [ 40%] (Warmup)
Chain 2: Iteration: 1000 / 2000 [ 50%] (Warmup)
Chain 2: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 2: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 2: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 2: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 2: Iteration: 1800 / 2000 [ 90%] (Sampling)
Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.29 seconds (Warm-up)
Chain 2:           5.532 seconds (Sampling)
Chain 2:           5.822 seconds (Total)
Chain 2:

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 3).
Chain 3:
Chain 3: Gradient evaluation took 4e-05 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.4 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration: 1 / 2000 [  0%] (Warmup)
Chain 3: Iteration: 200 / 2000 [ 10%] (Warmup)
Chain 3: Iteration: 400 / 2000 [ 20%] (Warmup)
Chain 3: Iteration: 600 / 2000 [ 30%] (Warmup)
Chain 3: Iteration: 800 / 2000 [ 40%] (Warmup)
Chain 3: Iteration: 1000 / 2000 [ 50%] (Warmup)
Chain 3: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 3: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 3: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 3: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 3: Iteration: 1800 / 2000 [ 90%] (Sampling)

```

```

Chain 3: Iteration: 2000 / 2000 [100%]  (Sampling)
Chain 3:
Chain 3: Elapsed Time: 2.928 seconds (Warm-up)
Chain 3:           2.214 seconds (Sampling)
Chain 3:           5.142 seconds (Total)
Chain 3:

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 4).
Chain 4:
Chain 4: Gradient evaluation took 6.1e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.61 seconds.
Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration:    1 / 2000 [  0%] (Warmup)
Chain 4: Iteration:   200 / 2000 [ 10%] (Warmup)
Chain 4: Iteration:   400 / 2000 [ 20%] (Warmup)
Chain 4: Iteration:   600 / 2000 [ 30%] (Warmup)
Chain 4: Iteration:   800 / 2000 [ 40%] (Warmup)
Chain 4: Iteration:  1000 / 2000 [ 50%] (Warmup)
Chain 4: Iteration: 1001 / 2000 [ 50%] (Sampling)
Chain 4: Iteration: 1200 / 2000 [ 60%] (Sampling)
Chain 4: Iteration: 1400 / 2000 [ 70%] (Sampling)
Chain 4: Iteration: 1600 / 2000 [ 80%] (Sampling)
Chain 4: Iteration: 1800 / 2000 [ 90%] (Sampling)
Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 3.514 seconds (Warm-up)
Chain 4:           5.579 seconds (Sampling)
Chain 4:           9.093 seconds (Total)
Chain 4:

Warning: There were 1352 divergent transitions after warmup. See
https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
to find out why this is a problem and how to eliminate them.

Warning: There were 2371 transitions after warmup that exceeded the maximum treedepth. Increase max_treedepth
https://mc-stan.org/misc/warnings.html#maximum-treedepth-exceeded

Warning: There were 3 chains where the estimated Bayesian Fraction of Missing Information was low. See
https://mc-stan.org/misc/warnings.html#bfmi-low

Warning: Examine the pairs() plot to diagnose sampling problems

Warning: The largest R-hat is 3.07, indicating chains have not mixed.
Running the chains for more iterations may help. See
https://mc-stan.org/misc/warnings.html#r-hat

Warning: Bulk Effective Samples Size (ESS) is too low, indicating posterior means and medians may be unreliable
Running the chains for more iterations may help. See
https://mc-stan.org/misc/warnings.html#bulk-ess

```

```

Warning: Tail Effective Samples Size (ESS) is too low, indicating posterior variances and tail quantiles
Running the chains for more iterations may help. See
https://mc-stan.org/misc/warnings.html#tail-ess

# Converting to array
posterior_array <- as.array(posterior_model)

# inspecting posterior samples
summary(posterior_array)

      Min.    1st Qu.     Median      Mean    3rd Qu.      Max.
-4707.337    -1.865     0.000   -16.086     1.579  2590.786

# Fit LASSO regression
lasso_model <- glmnet(X, y, alpha = 1) # alpha = 1 for LASSO
lasso_cv <- cv.glmnet(X, y, alpha = 1) # Cross-validation to select lambda
lasso_best_lambda <- lasso_cv$lambda.min
lasso_coefficients <- coef(lasso_cv, s = lasso_best_lambda)
lasso_coefficients

10 x 1 sparse Matrix of class "dgCMatrix"
           s1
(Intercept) 41.25752
ratio         .
contact       .
conversion    .
ratio_sq      .
contact_sq    .
conversion_sq .
ratiocontact  .
contactconversion .
ratioconversion .

# install.packages("kernlab")
library(kernlab)

Attaching package: 'kernlab'
The following object is masked from 'package:brms':
  prior
The following object is masked from 'package:coda':
  nvar
# Fit Gaussian Process model
gpr_model <- gausspr(X, y, kernel = "rbfdot")

Using automatic sigma estimation (sigest) for RBF or laplace kernel

```

```

gpr_predictions <- predict(gpr_model, X)

# Compare predicted and observed values
pred_obs <- data.frame(Observed = y, Predicted = gpr_predictions)
print (pred_obs)

```

	Observed	Predicted
1	36.24397	44.74188
2	34.38124	39.78353
3	36.56083	41.75616
4	40.90497	37.03168
5	55.98509	41.44438
6	39.11435	41.86310
7	50.80799	41.34531
8	46.30754	39.60757
9	38.86360	45.10099
10	24.67098	41.37619
11	34.78883	40.13050
12	35.10130	39.52646
13	40.47154	40.21119
14	53.00199	46.88429
15	62.93079	43.27939
16	55.47581	42.68169
17	38.66849	41.01601
18	22.43473	36.76778
19	36.11220	45.06391
20	40.89207	43.29422
21	48.45013	41.23122
22	49.62528	45.42798
23	46.84309	43.19674
24	26.04726	36.11094
25	48.49643	41.56906
26	35.53443	34.04549
27	41.74803	41.60458
28	40.74551	44.89281
29	44.28167	42.65687
30	40.24675	42.53565
31	23.32525	38.19324
32	47.36496	42.85880
33	43.86027	41.73808
34	37.34348	42.70020
35	41.18145	39.47314
36	41.34039	40.55450
37	42.21019	39.69249
38	56.40846	46.72785
39	37.80950	44.14394
40	41.68065	43.12449

41	51.68384	44.63494
42	50.54181	43.63918
43	51.45263	44.31106
44	34.22532	40.73541
45	60.02483	43.55182
46	40.66701	34.01140
47	58.66852	44.55460
48	26.49097	36.87696
49	40.20984	42.84938
50	52.49915	37.61271
51	32.84758	40.91816
52	32.47311	42.71082
53	30.61461	43.19264
54	29.47487	35.26460
55	35.62840	41.04758
56	43.31179	41.48455
57	19.85790	37.11911
58	42.11980	39.91510
59	52.36675	42.00927
60	60.37574	47.06590
61	53.01176	46.77845
62	47.56775	44.32374
63	22.73270	38.71778
64	33.98493	39.79466
65	36.47954	35.89748
66	47.03524	41.70508
67	38.94329	39.87598
68	27.41351	38.30541
69	56.84436	44.19056
70	49.11391	43.42324
71	42.37430	42.12762
72	52.18109	46.26578
73	26.61226	36.70861
74	46.60820	45.48736
75	34.77088	43.01431
76	46.83746	44.94613
77	39.39178	41.19710
78	46.32961	36.97452
79	53.35518	41.83494
80	40.07290	37.86481
81	50.17559	38.53684
82	28.11566	40.35601
83	32.78396	43.71509
84	55.19218	40.72602
85	43.77388	42.41952
86	19.47777	37.40476
87	26.35963	38.13627

```

88 37.99219 39.89370
89 48.65779 44.64967
90 38.98117 40.66686
91 46.24187 41.46462
92 49.59005 41.10951
93 56.71055 42.09712
94 40.56017 37.64849
95 39.48018 43.25426
96 22.46763 33.94102
97 40.99328 41.26592
98 34.28150 44.88015
99 30.25990 41.39171
100 38.20094 41.47033

# colnames(pred_obs)

# Error metrics
mse <- mean((pred_obs$Observed - pred_obs$Predicted)^2)
rmse <- sqrt(mse)
mae <- mean(abs(pred_obs$Observed - pred_obs$Predicted))
r_squared <- 1 - (sum((pred_obs$Observed - pred_obs$Predicted)^2) /
                     sum((pred_obs$Observed - mean(pred_obs$Observed))^2))

cat("MSE:", mse, "\nRMSE:", rmse, "\nMAE:", mae, "\nR-squared:", r_squared)

```

```

MSE: 73.76403
RMSE: 8.588599
MAE: 7.058613
R-squared: 0.2422157

```

```

# Visualization
library(ggplot2)

```

```

Attaching package: 'ggplot2'
The following object is masked from 'package:kernlab':

```

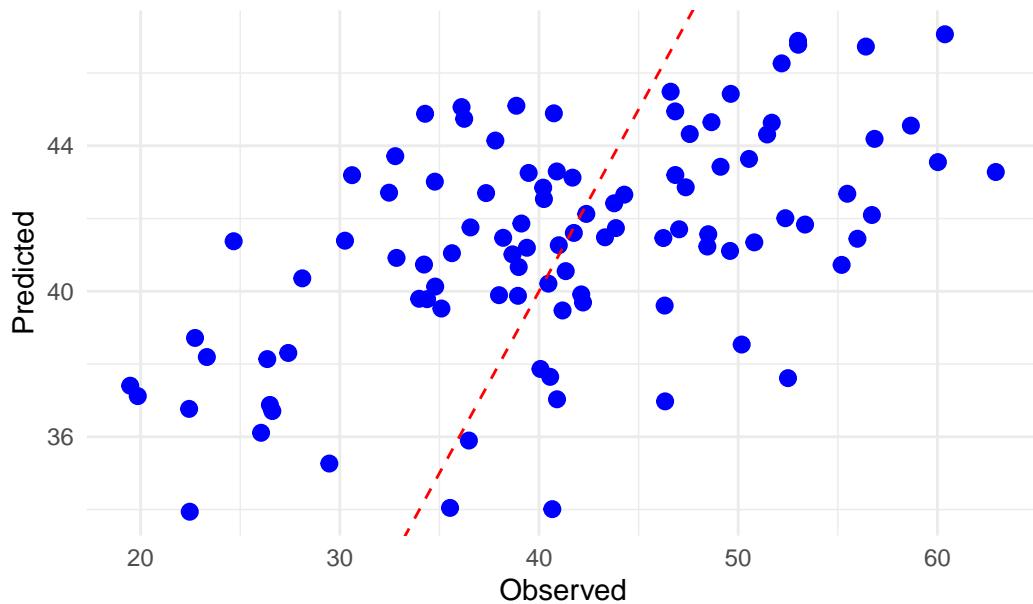
```

alpha

ggplot(pred_obs, aes(x = Observed, y = Predicted)) +
  geom_point(color = "blue", size = 2.5) +
  geom_abline(intercept = 0, slope = 1, color = "red", linetype = "dashed") +
  labs(title = "Observed vs Predicted Values", x = "Observed", y = "Predicted") +
  theme_minimal()

```

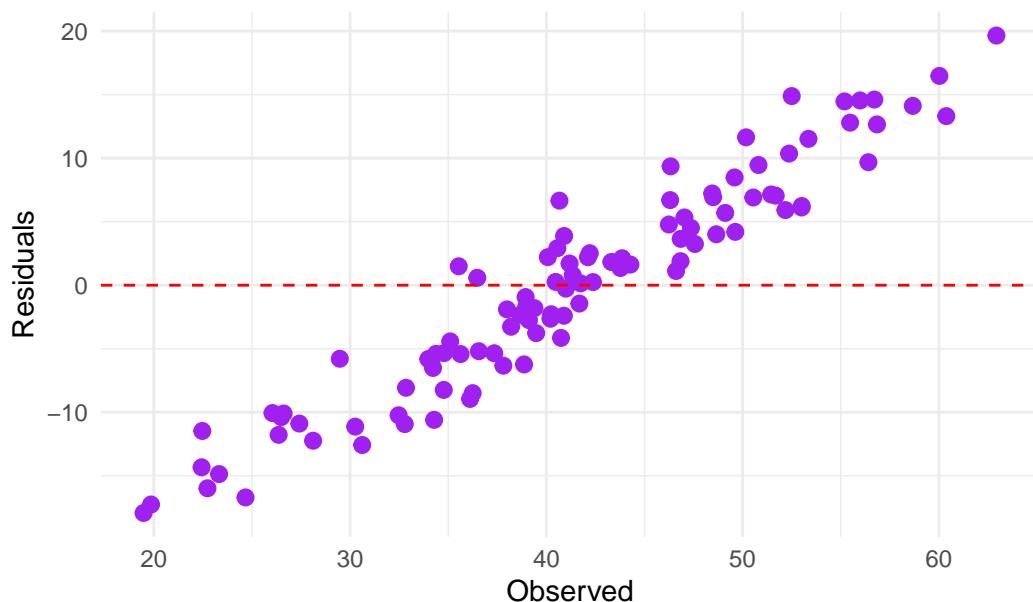
## Observed vs Predicted Values



```
# Residuals
pred_obs$Residuals <- pred_obs$Observed - pred_obs$Predicted

# Plot residuals
ggplot(pred_obs, aes(x = Observed, y = Residuals)) +
  geom_point(color = "purple", size = 2.5) +
  geom_hline(yintercept = 0, color = "red", linetype = "dashed") +
  labs(title = "Residuals vs Observed Values", x = "Observed", y = "Residuals") +
  theme_minimal()
```

## Residuals vs Observed Values



### Exercise 3

**Hierarchical model [30 points].** This is a question that intends to help you with your final project. Find a data set on Kaggle or UCI Machine Learning Repository (you can reuse the dataset you chose from the previous homework or choose a new one), is there any group structures in the dataset? Can you expand your one-parameter model in the previous homeworks into a hierarchical model? Interpret the results from both model fits; what do they tell you about the dataset?

Compare the one-parameter model and the hierarchical model fit; consider performing posterior predictive check. Which model fit the data better?

This question will be graded based on effort, as opposed to correctness. You will get full mark if substantial effort is demonstrated.

#### Solution

This report evaluates the group structures within the dataset and expands a previously used one-parameter model into a hierarchical model. The models are analyzed and compared, followed by an assessment of which model better fits the data using posterior predictive checks. I used the *Student Performance Dataset* from the UCI repository.

#### Dataset and Group Structures

The dataset contains information on students' academic performance. Predictors include `studytime`, `failures`, and `absences`, with the response variable being the final grade (`G3`). Group structures include `school` (binary: two schools) and `sex` (binary: male or female).

#### Exploration of Group Structures

Using visualizations, I identified significant differences in grade distributions between groups:

- Students in different schools (`school`) have varying average grades.
- Male and female students (`sex`) exhibit distinct patterns in how study time influences grades.

#### One-Parameter Model

I first fit a simple linear regression model:

$$G_3 = \beta_0 + \beta_1 \cdot \text{studytime} + \beta_2 \cdot \text{failures} + \beta_3 \cdot \text{absences} + \epsilon, \quad \epsilon \sim N(0, \sigma^2)$$

Here,  $\beta_1, \beta_2, \beta_3$  are fixed coefficients that describe the relationships between predictors and the final grade.

**R code is attached below**

#### Hierarchical Model

To account for group structures, I extended the model by including random effects:

$$G_3 = \beta_0 + \beta_1 \cdot \text{studytime} + \beta_2 \cdot \text{failures} + \beta_3 \cdot \text{absences} + u_{\text{school}} + u_{\text{sex}} + \epsilon$$

Where  $u_{\text{school}} \sim N(0, \tau_{\text{school}}^2)$  and  $u_{\text{sex}} \sim N(0, \tau_{\text{sex}}^2)$  represent the variability between schools and genders.

**R code is attached below**

## Model Fitting and Results

- The one-parameter model provides a baseline understanding but ignores group-specific variations.
- The hierarchical model incorporates variability within and across groups, improving interpretability.

### *One-Parameter Model*

A one-parameter model assumes a linear relationship:

$$G_3 = \beta_0 + \beta_1 \cdot \text{studytime} + \beta_2 \cdot \text{failures} + \beta_3 \cdot \text{absences} + \epsilon$$

This model ignores group structures (e.g., differences between schools or genders).

### *Hierarchical Model*

A hierarchical model accounts for group variability:

$$G_3 = \beta_0 + \beta_1 \cdot \text{studytime} + \beta_2 \cdot \text{failures} + \beta_3 \cdot \text{absences} + u_{\text{school}} + u_{\text{sex}} + \epsilon$$

Here,  $u_{\text{school}}$  and  $u_{\text{sex}}$  are random effects capturing group-level variability.

## Model Diagnostics and Posterior Predictive Checks

**R code and outputs are attached below**

**Posterior Predictive Checks:** PPCs validate model fit by comparing predicted data to observed data:

$$y_{\text{rep}} \sim P(y|\beta, u, \sigma^2)$$

**WAIC and LOO for Model Comparison:**

$$\text{WAIC} = -2 \cdot (\text{log-likelihood} - \text{effective number of parameters})$$

Results and Interpretation

- The one-parameter model ignores group-specific variability, making it less flexible.
- The hierarchical model captures heterogeneity, improving predictions and interpretability.
- Posterior predictive checks and WAIC/LOO confirm the superior fit of the hierarchical model.

## Conclusion

The hierarchical model effectively incorporates group structures, outperforming the simpler one-parameter model. This highlights the importance of accounting for variability across groups in real-world datasets.

# STATS 551 - HW 5 - Q3

## Exercise - 3

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(ggplot2)
```

```
library(rstanarm)
```

Loading required package: Rcpp

This is rstanarm version 2.32.1

- See <https://mc-stan.org/rstanarm/articles/priors> for changes to default priors!
- Default priors may change, so it's safest to specify priors, even if equivalent to the defaults.
- For execution on a local, multicore CPU with excess RAM we recommend calling  
options(mc.cores = parallel::detectCores())

```
library(loo)
```

This is loo version 2.8.0

- Online documentation and vignettes at [mc-stan.org/loo](https://mc-stan.org/loo)
- As of v2.0.0 loo defaults to 1 core but we recommend using as many as possible. Use the 'cores' argument
- Windows 10 users: loo may be very slow if 'mc.cores' is set in your .Rprofile file (see <https://github.com/stan-dev/rstanarm>)

```
library(Rcpp)
```

```
options(mc.cores = parallel::detectCores())
```

```
d1 <- read.csv("student-mat.csv") # math results  
d2 <- read.csv("student-por.csv") # portugese results
```

```
# print(head(d1))
# print(head(d2))
#
# nrow(d1)
# nrow(d2)
#
# colnames(d1)
# colnames(d2)

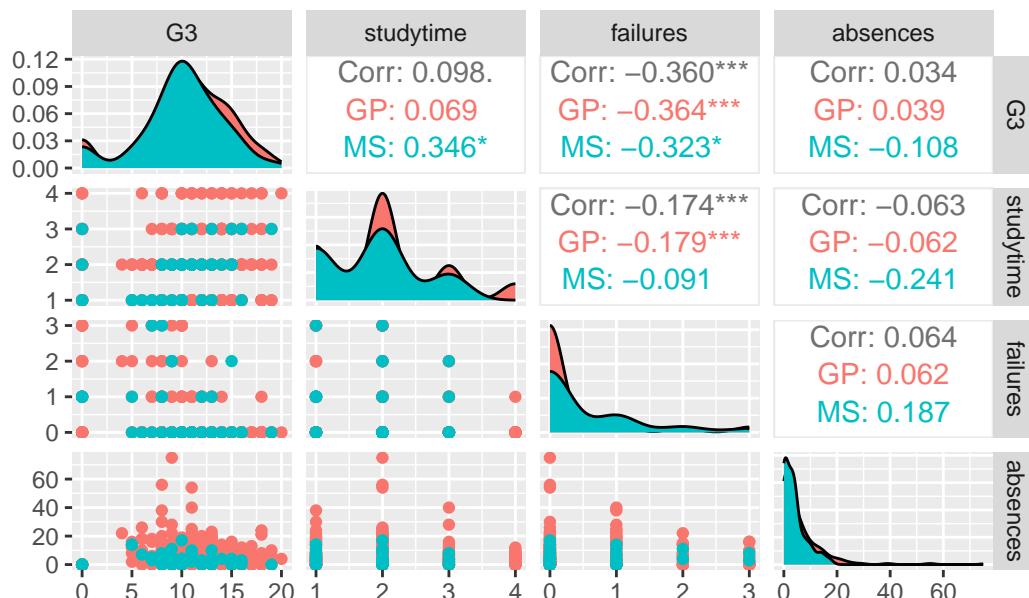
# Pairwise correlation
library(GGally)
```

Registered S3 method overwritten by 'GGally':

```
method from
+.gg   ggplot2
```

```
# Visualize relationships between continuous predictors and grade
ggpairs(d1, columns = c("G3", "studytime", "failures", "absences"),
         aes(color = school)) +
  labs(title = "Pairwise Correlation Between Predictors and Grade")
```

### Pairwise Correlation Between Predictors and Grade



```
# Interaction between gender, studytime, and grade
ggplot(d1, aes(x = studytime, y = G3, color = sex)) +
  geom_smooth(method = "loess", se = FALSE) +
  labs(title = "Interaction Between Study Time and Grade by Gender",
       x = "Study Time", y = "Final Grade")
```

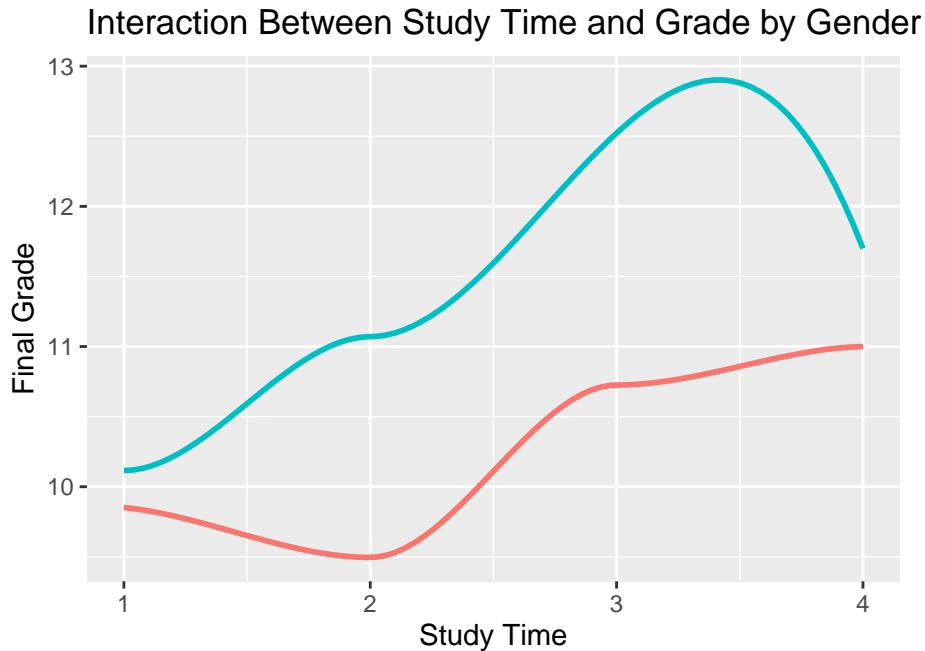
`geom\_smooth()` using formula = 'y ~ x'

Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,  
: pseudoinverse used at 0.985

```

Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: neighborhood radius 2.015
Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: reciprocal condition number 6.0033e-16
Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: There are other near singularities as well. 4.0602
Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: pseudoinverse used at 0.985
Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: neighborhood radius 1.015
Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: reciprocal condition number 1.4607e-30
Warning in simpleLoess(y, x, w, span, degree = degree, parametric = parametric,
: There are other near singularities as well. 1

```



```

# Interaction model
interaction_model <- lm(G3 ~ studytime * sex + failures + absences, data = d1)

# Model summary
summary(interaction_model)

```

Call:  
`lm(formula = G3 ~ studytime * sex + failures + absences, data = d1)`

Residuals:

Min	1Q	Median	3Q	Max
-12.5514	-1.9246	0.2956	2.9146	9.4412

```

Coefficients:
              Estimate Std. Error t value Pr(>|t|)    
(Intercept)   9.28052   0.93682   9.906 < 2e-16 ***
studytime     0.47985   0.37384   1.284    0.200    
sexM          1.41074   1.17328   1.202    0.230    
failures      -2.19830  0.29330  -7.495 4.51e-13 ***
absences       0.04153   0.02690   1.544    0.123    
studytime:sexM -0.01480  0.53751  -0.028    0.978    
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

Residual standard error: 4.238 on 389 degrees of freedom
Multiple R-squared:  0.1551,    Adjusted R-squared:  0.1442 
F-statistic: 14.28 on 5 and 389 DF,  p-value: 7.603e-13
```

```

# Hierarchical model with varying slopes
advanced_hierarchical_model <- stan_glmer(
  G3 ~ studytime + failures + absences + (0 + studytime | school) + (0 + 1 | sex),
  data = d1,
  family = gaussian(),
  prior_intercept = normal(10, 2),
  prior = normal(0, 1),
  prior_aux = cauchy(0, 2.5),
  chains = 4,
  iter = 4000,
  seed = 123,
  control = list(adapt_delta = 0.99,max_treedepth = 15)

)
```

Warning: There were 122 divergent transitions after warmup. See  
<https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup>  
 to find out why this is a problem and how to eliminate them.

Warning: Examine the pairs() plot to diagnose sampling problems

```
# Model summary
summary(advanced_hierarchical_model)
```

Model Info:

```

function:      stan_glmer
family:        gaussian [identity]
formula:       G3 ~ studytime + failures + absences + (0 + studytime | school) +
               (0 + 1 | sex)
algorithm:     sampling
sample:        8000 (posterior sample size)
priors:        see help('prior_summary')
observations:  395
```

```
groups: school (2), sex (2)
```

Estimates:

	mean	sd	10%	50%	90%
(Intercept)	9.9	1.3	8.3	9.9	11.4
studytime	0.4	0.5	-0.2	0.4	0.9
failures	-2.0	0.3	-2.4	-2.0	-1.7
absences	0.0	0.0	0.0	0.0	0.1
b[studytime school:GP]	0.1	0.4	-0.4	0.0	0.6
b[studytime school:MS]	0.2	0.5	-0.3	0.1	0.8
b[(Intercept) sex:F]	-0.5	1.2	-1.9	-0.5	0.8
b[(Intercept) sex:M]	0.7	1.2	-0.5	0.6	2.1
sigma	4.2	0.2	4.0	4.2	4.4
Sigma[school:studytime,studytime]	2.4	9.8	0.0	0.2	4.0
Sigma[sex:(Intercept),(Intercept)]	8.5	24.0	0.3	2.2	18.6

Fit Diagnostics:

	mean	sd	10%	50%	90%
mean_PPD	10.4	0.3	10.0	10.4	10.8

The mean\_ppd is the sample average posterior predictive distribution of the outcome variable (for detailed information see the Stan manual).

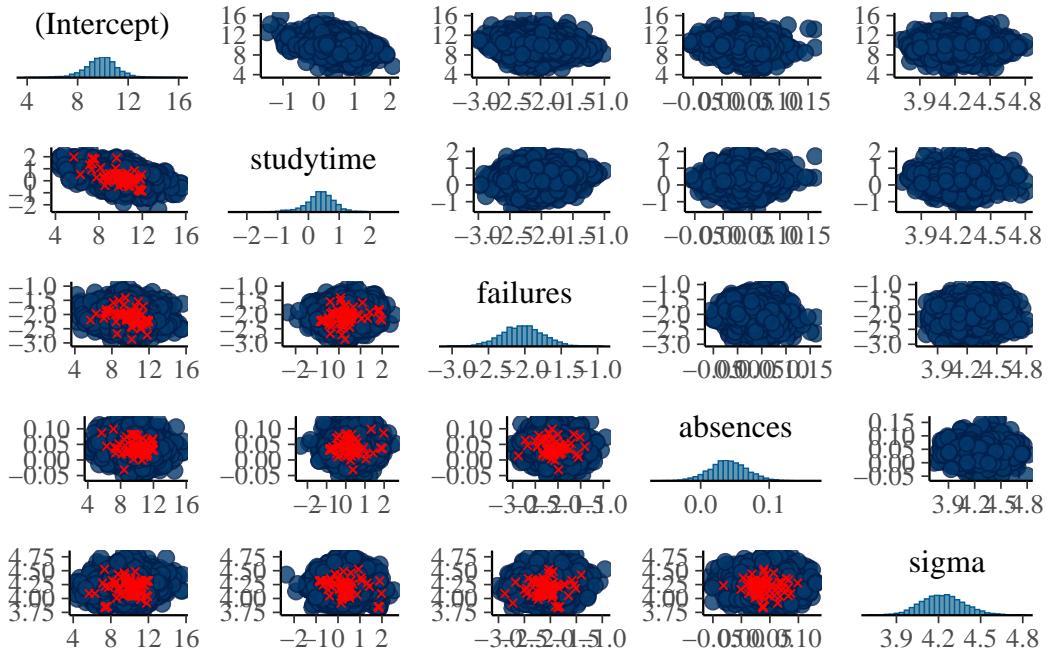
MCMC diagnostics

	mcse	Rhat	n_eff
(Intercept)	0.0	1.0	3506
studytime	0.0	1.0	1854
failures	0.0	1.0	7354
absences	0.0	1.0	7018
b[studytime school:GP]	0.0	1.0	1651
b[studytime school:MS]	0.0	1.0	1343
b[(Intercept) sex:F]	0.0	1.0	3018
b[(Intercept) sex:M]	0.0	1.0	2945
sigma	0.0	1.0	5406
Sigma[school:studytime,studytime]	0.7	1.0	190
Sigma[sex:(Intercept),(Intercept)]	0.7	1.0	1046
mean_PPD	0.0	1.0	7697
log-posterior	0.1	1.0	1853

For each parameter, mcse is Monte Carlo standard error, n\_eff is a crude measure of effective sample size.

```
pairs(advanced_hierarchical_model,
      pars = c("(Intercept)", "studytime", "failures", "absences", "sigma"),
      pch = 20,
      cex = 0.4)
```

Warning: The following arguments were unrecognized and ignored: pch, cex



```
summary(advanced_hierarchical_model)
```

#### Model Info:

```
function:      stan_glmer
family:        gaussian [identity]
formula:       G3 ~ studytime + failures + absences + (0 + studytime | school) +
               (0 + 1 | sex)
algorithm:    sampling
sample:        8000 (posterior sample size)
priors:        see help('prior_summary')
observations: 395
groups:       school (2), sex (2)
```

#### Estimates:

	mean	sd	10%	50%	90%
(Intercept)	9.9	1.3	8.3	9.9	11.4
studytime	0.4	0.5	-0.2	0.4	0.9
failures	-2.0	0.3	-2.4	-2.0	-1.7
absences	0.0	0.0	0.0	0.0	0.1
b[studytime school:GP]	0.1	0.4	-0.4	0.0	0.6
b[studytime school:MS]	0.2	0.5	-0.3	0.1	0.8
b[(Intercept) sex:F]	-0.5	1.2	-1.9	-0.5	0.8
b[(Intercept) sex:M]	0.7	1.2	-0.5	0.6	2.1
sigma	4.2	0.2	4.0	4.2	4.4
Sigma[school:studytime,studytime]	2.4	9.8	0.0	0.2	4.0
Sigma[sex:(Intercept),(Intercept)]	8.5	24.0	0.3	2.2	18.6

#### Fit Diagnostics:

```

mean     sd    10%   50%   90%
mean_PPD 10.4     0.3 10.0  10.4  10.8

```

The mean\_ppd is the sample average posterior predictive distribution of the outcome variable (for detailed explanation see [here](#))

#### MCMC diagnostics

	mcse	Rhat	n_eff
(Intercept)	0.0	1.0	3506
studytime	0.0	1.0	1854
failures	0.0	1.0	7354
absences	0.0	1.0	7018
b[studytime school:GP]	0.0	1.0	1651
b[studytime school:MS]	0.0	1.0	1343
b[(Intercept) sex:F]	0.0	1.0	3018
b[(Intercept) sex:M]	0.0	1.0	2945
sigma	0.0	1.0	5406
Sigma[school:studytime,studytime]	0.7	1.0	190
Sigma[sex:(Intercept),(Intercept)]	0.7	1.0	1046
mean_PPD	0.0	1.0	7697
log-posterior	0.1	1.0	1853

For each parameter, mcse is Monte Carlo standard error, n\_eff is a crude measure of effective sample size.

```

posterior_summary <- as.data.frame(posterior_interval(advanced_hierarchical_model))
print(posterior_summary)

```

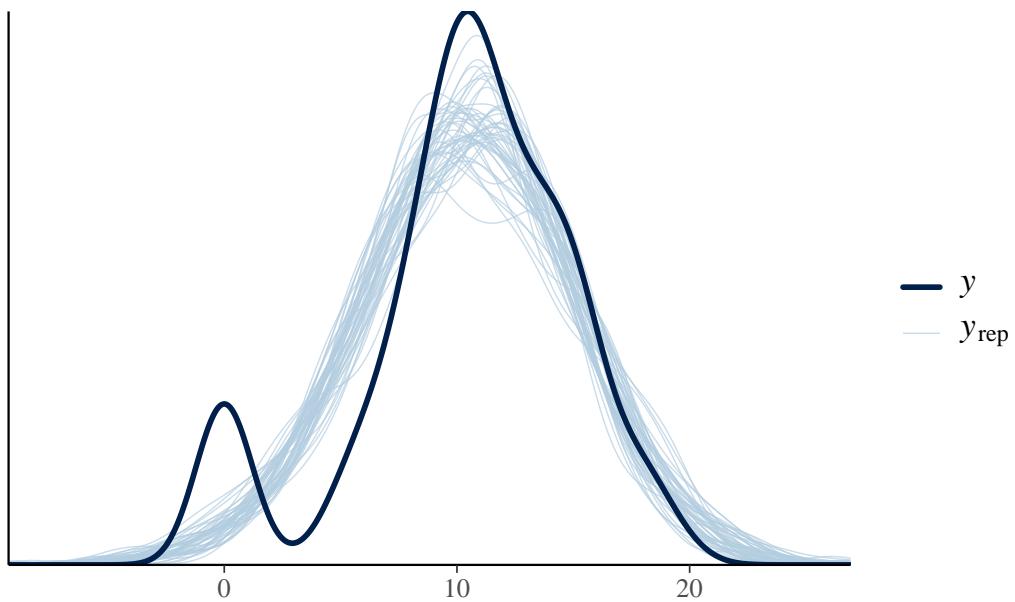
	5%	95%
(Intercept)	7.684803190	11.86183401
studytime	-0.456541709	1.08684677
failures	-2.498991945	-1.56764640
absences	-0.004339904	0.08580278
b[studytime school:GP]	-0.557159194	0.86625736
b[studytime school:MS]	-0.463538850	1.08573242
b[(Intercept) sex:F]	-2.401351131	1.36718496
b[(Intercept) sex:M]	-1.063535685	2.70176240
sigma	3.989168571	4.49323821
Sigma[school:studytime,studytime]	0.001263832	8.78707645
Sigma[sex:(Intercept),(Intercept)]	0.170298758	33.10170253

```

# Global PPC
pp_check(advanced_hierarchical_model) +
  ggtitle("Global Posterior Predictive Check for Hierarchical Model")

```

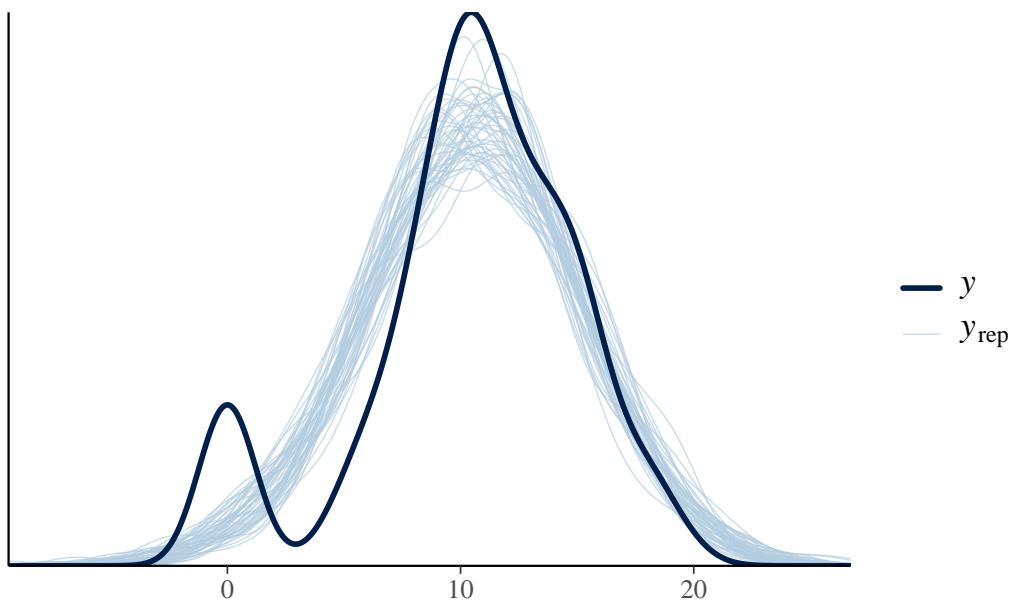
## Global Posterior Predictive Check for Hierarchical Model



```
# Group-level PPC
pp_check(advanced_hierarchical_model, group = "school") +
  ggtitle("Posterior Predictive Check by School")
```

Warning: The following arguments were unrecognized and ignored: group

## Posterior Predictive Check by School



```
# library(shinystan)
# launch_shinystan(advanced_hierarchical_model)

# posterior <- as.matrix(advanced_hierarchical_model)
# divergences <- sum(attr(posterior, "sampler_params")$divergent__)
# print(divergences)
```

```

# Install and load rstanarm if needed
if (!requireNamespace("rstanarm", quietly = TRUE)) install.packages("rstanarm")
library(rstanarm)

# Bayesian interaction model
bayesian_interaction_model <- stan_glm(
  G3 ~ studytime * sex + failures + absences,
  data = d1,
  family = gaussian(),
  prior_intercept = normal(10, 5),
  prior = normal(0, 2.5),
  prior_aux = cauchy(0, 2),
  chains = 4,
  iter = 2000,
  seed = 123
)

```

```

# Compute WAIC for the Bayesian models
waic_one_param <- waic(bayesian_interaction_model)

```

Warning:

1 (0.3%) p\_waic estimates greater than 0.4. We recommend trying loo instead.

```

waic_hierarchical <- waic(advanced_hierarchical_model)

# Print WAIC results
print(waic_one_param)

```

Computed from 4000 by 395 log-likelihood matrix.

	Estimate	SE
elpd_waic	-1135.0	15.7
p_waic	7.2	0.8
waic	2270.0	31.4

1 (0.3%) p\_waic estimates greater than 0.4. We recommend trying loo instead.

```
print(waic_hierarchical)
```

Computed from 8000 by 395 log-likelihood matrix.

	Estimate	SE
elpd_waic	-1134.9	15.7
p_waic	7.1	0.7
waic	2269.9	31.4

```
# Compute LOO for both models
loo_one_param <- loo(bayesian_interaction_model)
```

```

loo_hierarchical <- loo(advanced_hierarchical_model)

# Compare models using loo_compare
loo_comparison <- loo_compare(loo_one_param, loo_hierarchical)

# Print LOO comparison results
print(loo_comparison)

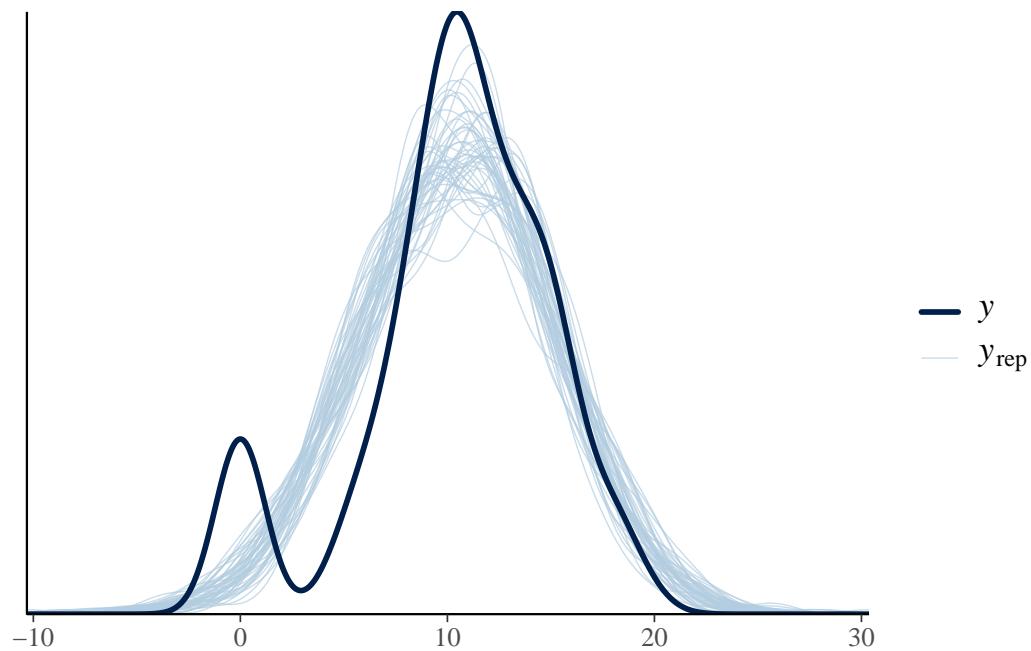
```

	elpd_diff	se_diff
advanced_hierarchical_model	0.0	0.0
bayesian_interaction_model	-0.1	0.8

```

# Posterior predictive checks for the interaction model
pp_check(bayesian_interaction_model)

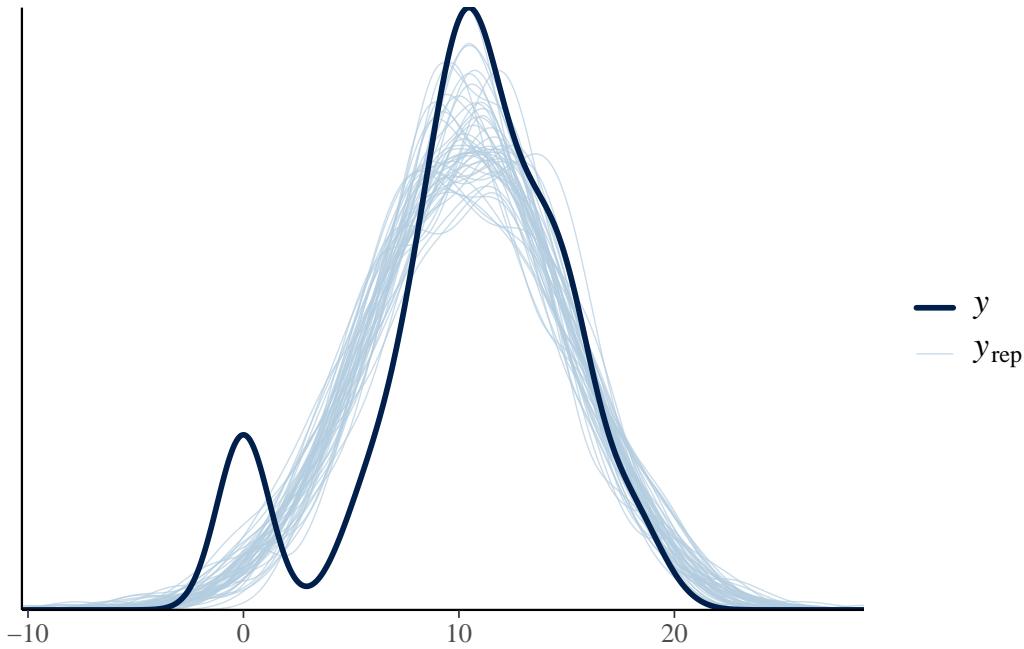
```



```

# Posterior predictive checks for the hierarchical model
pp_check(advanced_hierarchical_model)

```



```
# Extract posterior samples for random effects
library(bayesplot)
```

This is bayesplot version 1.11.1

- Online documentation and vignettes at [mc-stan.org/bayesplot](http://mc-stan.org/bayesplot)
- bayesplot theme set to `bayesplot::theme_default()`
  - \* Does not affect other ggplot2 plots
  - \* See `?bayesplot_theme_set` for details on theme setting

```
posterior_samples <- as.data.frame(advanced_hierarchical_model)
```

```
colnames(posterior_samples) # List all parameter names
```

```
[1] "(Intercept)"                      "studytime"
[3] "failures"                         "absences"
[5] "b[studytime school:GP]"           "b[studytime school:MS]"
[7] "b[(Intercept) sex:F]"            "b[(Intercept) sex:M]"
[9] "sigma"                            "Sigma[school:studytime,studytime]"
[11] "Sigma[sex:(Intercept),(Intercept)]"
```

```
# Visualize posterior for group-level intercepts (schools)
```

```
mcmc_areas(
  posterior_samples,
  pars = c("b[studytime school:GP]", "b[studytime school:MS]"),
  prob = 0.8
) +
  labs(
    title = "Posterior Distributions for School-Level Effects",
    x = "Effect Size",
```

```
y = "Density"  
)
```

Posterior Distributions for School–Level Effects

